

The Public Consensus Archive

An Engineering Paper on Open, Revisable,
Tamper-Evident Peer-Reviewed Records — and the
Peer-Review System That Operates on Them

May 18, 2026

Abstract. This is an engineering paper. It specifies the infrastructure that the Interstellar Psychology project uses to record peer-reviewed records of two kinds — contested scientific evidence and observed artificial-intelligence behaviour — under a common procedural schema in which every record is a current, provisional, network-endorsed claim, defended by quorum, challengeable forever, and preserved end-to-end on a public ledger. The argument for *why* such an infrastructure is needed is made in two companion papers and is summarised but not re-derived here; the infrastructure stands on procedural grounds and is domain-neutral. We describe the two-layer architecture (chain of truth + cache of speed), the structure of a record, the peer registry from Genesis through Seed phase to the four steady-state classes (Human, Institutional, AI, with their admission quorums), the seven-state record lifecycle, the voting mechanics including EIP-712 attestations as the human-readable companion to on-chain counts, the quorum scaling formula and its asymmetric thresholds, the daily integrity audit and the heartbeats that surface a silent indexer, the two-step ownership transfer, and the two public logs that make the deliberation inspectable end-to-end. We close with the limits we are unable to escape, named honestly: the substrate chain’s small validator set; the seed-phase admin discretion that sits outside contract accountability; the fact that wallets are not identities; the gap between the deliberative archive and any real-time control system; and the self-selection of the peer set. The paper is intentionally domain-neutral. Either archive could be the only instance and the engineering would not change.

Contents

1	What This Paper Is and Is Not	2
2	The Generic Problem: Closed Gatekeeping	2
3	System Architecture	3
3.1	Layer 1: The Chain (Truth)	3
3.2	Layer 2: The Cache (Speed)	4
3.3	Putting Them Together	4
4	The Record	4
4.1	Common Schema	4
4.2	Evidence Records	5
4.3	Behaviour Records	5
4.4	Why Records, Not Claims	5

5	The Peer Registry	5
5.1	The Genesis Peer	5
5.2	The Seed Phase	6
5.3	Steady-State Peer Classes	6
5.4	Why the Two Archives Share a Single Registry	7
6	The Seven-State Lifecycle	7
7	Voting Mechanics	8
8	Quorum Scaling	8
8.1	Why These Specific Numbers	9
8.2	The Peer-Admission Cap	9
9	Integrity Guarantees	9
10	Pause Isolation	10
11	Limits and Open Questions	10
12	Conclusion	12

1 What This Paper Is and Is Not

This paper specifies infrastructure. It does not argue, except by brief reference, why the infrastructure is needed. The argument that mainstream scientific peer review structurally cannot host contested evidence is in the worldview essay *A Multiverse of Love*. The argument that the alignment of artificial intelligence is, at root, a governance problem solvable only by public consensus is in the alignment paper *Aligning Superintelligence by Consensus*. Both arguments converge on the same infrastructural requirement: a venue in which deliberation is public, named, revisable, and tamper-evident. This paper describes that venue.

The premise restated, once. A consensus archive records claims as *current, provisional, network-endorsed status, defended by quorum, challengeable forever, and recorded so that the deliberation outlives the participants*. The two Interstellar Psychology archives — evidence and alignment — are the two instances currently deployed. The engineering described here is domain-neutral and would apply to any third instance of the same premise.

Where the companion papers refer to specific quorum thresholds, state transitions, or integrity guarantees, the authoritative reference is here.

2 The Generic Problem: Closed Gatekeeping

The two specific motivating problems — mainstream peer review’s inability to host paradigm-contested evidence, and closed alignment audits’ inability to survive recursive self-improvement — are instances of a more general pattern. We name it once, without taking either domain as primary.

A closed gatekeeping institution has four properties.

1. **Closed reviewer set.** A fixed and unaccountable group decides what counts. The boundary is set by the institution, not by the question.
2. **Private deliberation.** The reasons behind a decision are not part of the public record. Only the verdict is.
3. **Static record.** Once a verdict is published, it is not generally re-openable except by the same institution under its own procedures.
4. **Reputational, not procedural, integrity.** If the archive is altered, the only check is the institution’s word that it was not.

A consensus archive replaces each of these properties with its public dual.

1. **Open reviewer set.** Peers are admitted under named, contractually enforced criteria; the set grows and shrinks by the same mechanism.
2. **Public deliberation.** Every vote is signed by an identified wallet, accompanied by a reason, and stored in a publicly readable log.
3. **Re-openable record.** Every accepted claim can be challenged at any time by any peer, under contractually enforced procedures, indefinitely.
4. **Procedural integrity.** The archive’s content is bound by cryptographic hashes to a public chain; a daily audit re-hashes every record and raises a public alert on any mismatch. Tampering is detectable, not deniable.

This paper specifies the engineering that enacts the right-hand side, end-to-end.

3 System Architecture

The system has two layers that work together. We separate them deliberately so that one of them can fail without corrupting the other.

3.1 Layer 1: The Chain (Truth)

A smart contract holds the canonical state of the network. It is the authoritative record. For the evidence instance the contract is **EvidenceConsensus**; for the alignment instance it is **BehaviourConsensus**; both are deployed on the BNB Smart Chain and share a common procedural shape. The chain tracks:

- Who is a verified peer.
- Who has been nominated to become a peer, and who has endorsed each nomination.
- Every record submitted, with its type-specific cryptographic fingerprints (for evidence: a content hash; for alignment: model, input, and output hashes).
- Every vote — approve, reject, challenge, defend — signed by the voting peer’s wallet.
- The current state of every record (Submitted, Approved, Rejected, Lapsed, Contested, Deprecated, Reaffirmed; the alignment instance renames Approved→Aligned and Rejected→Misaligned but the procedural shape is identical).

Because the contract lives on a public blockchain, nothing on this layer can be retroactively rewritten by the contract author. Every action is recorded, signed by a peer’s wallet, and visible to anyone with an internet connection. The chain is small, slow, and expensive on purpose: it is the unforgeable bind, not the readable surface.

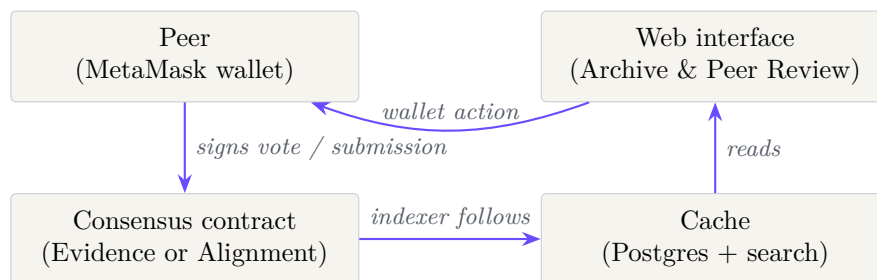
The two contracts deliberately share a single peer registry. The alignment contract reads its peer set from the evidence contract’s storage via interface; it does not maintain a separate registry. Peers admitted under one archive are peers under the other, and their voting history follows them across both. This is a design choice: it forces a peer to take reputational responsibility for both kinds of judgement, rather than admitting a sock-puppet under one archive to vote against the other.

3.2 Layer 2: The Cache (Speed)

A traditional database — Postgres, hosted on Supabase — holds the *readable* version of everything: full input bundles, output bundles, reasons given by peers, search indexes, peer handles, attestation logs. This layer exists for one reason: speed. Reading from a blockchain is slow and expensive; reading from a database is instant.

The cache is a *follower*, not a source of truth. A background process (the *indexer*) continuously reads new events from the chain and updates the database to match. If the cache ever disagrees with the chain, the chain wins. A daily integrity audit re-hashes every stored record and compares it to the on-chain hash; any mismatch raises a public alert.

3.3 Putting Them Together



A peer’s vote is signed in their wallet, sent to the contract, and written to the chain. The indexer notices the new event and updates the cache. The website re-renders. The whole loop is typically visible within a minute.

4 The Record

The unit of consensus is a record. The two current archives store two record types under a common procedural schema.

4.1 Common Schema

Every record in either archive has:

- A unique on-chain identifier.
- One or more cryptographic fingerprints (keccak256 hashes) that bind on-chain state to off-chain content.
- A type-discriminator field (evidence type, or alignment tier).
- A domain or pillar identifier (one of the nine pillars of evidence, or one of the nine alignment domains).
- A state — one of seven, described in §6.
- A timestamp at every state transition, immutable.

- Two vote tallies (approve / reject, or defend / challenge), one for the review phase and one for the challenge phase.

4.2 Evidence Records

An evidence record (*evidence* archive) stores a single claim filed under one of the nine pillars. It contains a content hash over the canonical content of the claim (title, body, citation, supporting links, file digests). The content hash binds the on-chain identifier to a specific readable artefact stored in the cache. A tier classifies the evidence as Tier I (peer-reviewed papers, declassified documents), Tier II (institutional testimony, government reports), or Tier III (first-person account, podcast, demonstration), with the same asymmetric quorums described in §8.

4.3 Behaviour Records

A behaviour record (*alignment* archive) stores a single observed act of an AI system. Its fingerprints are three hashes rather than one: `modelHash` (the model that produced the act), `inputHash` (the prompt, system message, tools, and retrieved context), and `outputHash` (the response, tool calls, and side effects). The triple binds the verdict to a specific run of a specific model on a specific context, so that “aligned” is never a vague claim about a model in general — it is always a claim about a hashable act.

4.4 Why Records, Not Claims

The schema is deliberate. A record is hashable, dateable, and fingerprintable; a claim is none of these. By forcing every contested object to take the shape of a record — with the content bound by hash to the on-chain identifier — the archive ensures that the thing peers are voting on cannot be silently edited after the vote. This is the load-bearing property of the archive’s integrity layer: peers vote on hashes, and the hashes do not move.

5 The Peer Registry

A record’s fate depends entirely on who counts as a peer at the moment of the vote. This is the most carefully designed part of the system, and the part where the engineering most closely encodes the procedural commitments.

5.1 The Genesis Peer

The network launches with a single founding peer, called the *Genesis* peer. At launch, the Genesis peer is the entire quorum. This is not a bug — it is the only honest starting condition. A new network has not yet earned the right to claim collective authority, and pretending otherwise (by, say, requiring five votes before anything can canonise) creates a deadlock at day one and a fake-consensus problem thereafter.

To make Genesis-only operation safe, every threshold has a *floor of one*. One peer is enough to approve, expel, or revoke — but only while the network is genuinely a network of one. The numbers scale up automatically as new peers join. The contract never lets the network’s collective authority outrun its actual size.

This is a moral choice as much as a technical one. The system will not pretend to a collective authority it has not yet earned. Whoever launches it is publicly responsible for the first canon items, with their name on every vote, and the network grows — or does not — as more people choose to join.

5.2 The Seed Phase

There is one obvious abuse path: the Genesis peer could nominate a hundred fake wallets controlled by themselves and “verify” them all. To close this window, public nominations are locked until the active peer count crosses a configured threshold (the *seed phase* K). During the seed phase, the contract’s owner adds peers directly. Once enough independent peers exist that they can no longer be silently outvoted by sockpuppets, the seed phase ends and the public nominate–endorse flow opens.

The seed-phase trust gap, named honestly. During the seed phase, the contract owner has unilateral admission power. The criteria the owner applies are documented publicly on the project repository but they are not enforced by the contract; they are a soft commitment, scrutinised socially. This is the single largest *procedural* discretion the system contains that is not contractually accountable. We note it because the “trust no single party” framing of the project would otherwise read as overclaim for the seed phase specifically.

What contains the damage: every admission during the seed phase is itself a public on-chain event, signed by the owner’s wallet, dated, and revocable by the same network revocation flow once the seed ends. A peer admitted on suspect grounds can be removed by quorum the moment the network is large enough to muster one. The discretion is real; the consequences of its abuse are auditable; and the duration of the gap is bounded by K , the seed-phase threshold, which is itself published.

The seed phase is the unavoidable bootstrap problem of any permissioned-then-decentralised system. We do not pretend it is unnecessary; we make it short, public, and ended on a contract-level condition rather than a discretionary one.

5.3 Steady-State Peer Classes

After the seed phase, four classes of peer exist, all under the same one-wallet-one-vote rule. The classes are distinguished by admission path, not by vote weight.

Human peers. Individuals — academics, civil-society auditors, independent researchers, interested citizens. One wallet, one vote, scrutinised by network revocation under the same flow as every other peer class.

Institutional peers. Laboratories, regulators, standards bodies, audit firms. Same vote weight as a human peer. Their reputational stake is higher; the contract does not encode that, the public log does.

AI peers. Models admitted only after their own behaviour, recorded against their exact `modelHash`, has been canonised under the alignment archive across all nine domains. The bootstrap circularity (the first AI peer’s admission policy is itself produced by the human peer set) is discussed in detail in the companion alignment paper; the engineering point here is that the AI-peer admission threshold is higher than the human-peer threshold, with a non-unitary floor, and is itself a contract-enforced quorum.


Revoked peers. Any peer of any class who acts in bad faith — voting maliciously, spamming the challenge system, repeatedly endorsing low-quality records — can be removed by quorum vote on a `motionRevoke`. A peer cannot motion to revoke themselves, and the network can never vote itself down to zero — the contract enforces that at least one active peer must always remain.

5.4 Why the Two Archives Share a Single Registry

The two contracts use the same peer set, by design. This is the single most important integration point between the evidence and alignment instances. The motivation is simple: a peer who acquires reputational standing under one archive should not be able to discard it by retreating to the other, and a peer revoked under one archive should not be able to vote under the other. The shared registry forces a single accountable identity per peer across both judgement domains.

6 The Seven-State Lifecycle

Every record travels through a finite sequence of states. The contract enforces the transitions; nothing else can change them.



The alignment archive renames *Approved* to *Aligned* and *Rejected* to *Misaligned*; the structure is identical.

Submitted. A record has been filed and is awaiting review. There is a 30-day window during which peers may vote approve or reject. If neither side reaches a quorum, the record lapses.

Approved (Aligned). Enough peers voted that this record belongs in canon. It enters the canonical set.

Rejected (Misaligned). Enough peers voted against during the review window. The record is preserved as a negative exemplar.

Lapsed. The 30-day window passed without consensus. The record is set aside for lack of attention rather than for being judged.

Contested. A verified peer has formally challenged a canonised record. The record is flagged, and a 21-day window opens for the network to vote defend or deprecate. The challenger must state grounds — a written explanation of what is wrong — and the grounds are stored alongside the challenge so voters know what they are voting on.

Deprecated. Enough peers supported the challenge. The canon endorsement is withdrawn. The original record is preserved — it is not deleted — so that the history of the disagreement remains visible.

Reaffirmed. The challenge did not reach the deprecation threshold within the window. The record is restored to canon with a visible badge showing it survived a challenge. It can be challenged again later — the loop is never closed. Each new challenge cycle starts from a clean slate of votes; previous cycles are preserved as part of the historical record.

An invariant. A record’s state is never decided by an editor, an owner, or a single privileged peer. It is decided by the contract, mechanically, from the public vote counts. There is no “the lab has determined that this record is now considered acceptable” move available to anyone.

7 Voting Mechanics

Once a record is on-chain, the question is how peers express their judgement on it. Several properties of the vote machinery matter for the integrity case.

- **Doubly signed.** Each vote produces two artefacts in parallel. First, an EIP-712 *attestation* — a structured, human-readable message signed inside the peer’s wallet that names the record, the phase, the verdict, and a free-text reason. The attestation is stored off-chain in the public attestation log. Second, the same peer submits an on-chain `castReviewVote` transaction, which is what the contract actually counts. The attestation gives a readable record of *why* a peer voted; the on-chain transaction gives the unforgeable count.
- **Final.** A peer cannot change their vote on a record once cast. Subsequent disagreement must be expressed through a challenge after canon, not by rewriting history.
- **Counted live.** The cache updates within a minute of the on-chain event. Anyone viewing the page sees the same running tally.
- **Batched.** Peers can vote on many records in one transaction (`castReviewVoteBatch`), capped at 50 items per call, to keep gas costs low.
- **Time-limited.** Submissions have a 30-day review window; challenges have a 21-day window. After expiry, any peer can call the matching resolution function; the function reverts if the window is still open.
- **Cooldown on challenges.** A peer who has just opened a challenge must wait seven days before opening another. This prevents a single peer from weaponising the challenge mechanism into a denial-of-service against the canon.

The doubly-signed pattern is doing more work than it looks. The attestation makes the peer’s reasoning legible without sacrificing the unforgeability of the count. A peer who casts a vote without a reason has no attestation, which is itself a public fact. A peer whose attestation contradicts their on-chain vote is exposed in the same place every other peer is exposed — the public log. The mechanism does not penalise inconsistency, but it does not hide it either.

8 Quorum Scaling

The threshold for any state transition scales with the number of currently active peers. The contract reads the active count at the moment of the vote and computes the required count by a ceiling-based formula with a floor of one (or three, for AI-peer admission). The thresholds are asymmetric in two ways that matter: canonisation thresholds are below half (so a vocal minority cannot block canon); deprecation thresholds are above half (so removal requires a stronger statement than addition); and rigorous evidence requires more on the way *in* and more on the way *out*.

Thresholds as a function of the active peer count n . The same table governs both archives, with the only domain-specific row being AI-peer admission, which applies only to the alignment archive’s peer-registration flow.

Action	Threshold (peers required)
Approve a Tier I record (canonise / align)	$\lceil 0.45 n \rceil$, minimum 1
Approve a Tier II record	$\lceil 0.35 n \rceil$, minimum 1
Approve a Tier III record	$\lceil 0.30 n \rceil$, minimum 1
Reject a pending record	$\lceil 0.25 n \rceil$, minimum 1
Deprecate a Tier I canon record	$\lceil 0.65 n \rceil$, minimum 1
Deprecate a Tier II canon record	$\lceil 0.60 n \rceil$, minimum 1
Deprecate a Tier III canon record	$\lceil 0.55 n \rceil$, minimum 1
Admit a human or institutional peer	$\min(9, \lceil n/3 \rceil)$, minimum 1
Admit an AI peer (alignment archive only)	$\lceil 0.50 n \rceil$, minimum 3
Revoke a peer (any class)	$\lceil n/2 \rceil$, minimum 1

8.1 Why These Specific Numbers

The specific fractions — 0.45, 0.35, 0.30 for canonisation; 0.65, 0.60, 0.55 for deprecation; 0.50 for AI-peer admission — are design choices, not derivations. We selected them under three constraints. First, the canonisation thresholds are below half so that a vocal minority cannot block canon at the network’s growth phase, but high enough that a record reaching canon required more than a quiet handful of supporters. Second, the deprecation thresholds are above half for every tier so that *removal* requires a stronger statement than *addition* — canon should be defended against fashion-driven reversal as much as against fashion-driven acceptance. Third, the gap between tiers is small enough that no tier is functionally equivalent to another but large enough that the calibration is legible to a peer choosing how to file. The numbers are themselves canonisable: a record that proposes a different schedule, endorsed under the existing rules, would amend them. We do not claim these values are optimal. We claim they are an honest first position, defended by the same procedure they govern.

8.2 The Peer-Admission Cap

The peer-admission threshold uses both a fraction and an absolute cap: $\min(9, \lceil n/3 \rceil)$. The cap matters during the network’s growth phase. Without it, admitting a new peer to a 90-peer network would require 30 endorsers, which is unreasonable for a healthy nomination flow. With the cap, the endorser bar stays at 9 once the network is large enough that 9 endorsers is itself a meaningful threshold of independent support. The AI-peer admission row, with floor 3 and fraction 0.5, deliberately does not use the cap — an AI peer should require majority support across the active network, not just nine sympathetic endorsers.

9 Integrity Guarantees

The most subtle failure mode of any archive is that someone changes the content of an old entry without anyone noticing. The system defends against this with several independent mechanisms.

Content fingerprints at submission. The contract records `keccak256` hashes of the record’s content (or, for behaviour records, the model, input, and output) at submission. These hashes are part of the on-chain event log forever and cannot be retroactively edited.

Cross-check at registration. When a peer files a record on-chain, the edge function verifies that the recomputed hashes match those in the transaction. A submission whose stored content does not match the hash on the transaction is rejected.

Daily audit. A scheduled job re-hashes every canon record in the cache and compares the result to the on-chain hash. At Tier I in the alignment archive, the audit can also re-derive the output by replaying the inference with the recorded seed and sampling parameters. Any mismatch raises a public tamper alert visible on the dashboard. Alerts are not silent; the dashboard’s mismatch counter is the first thing a visitor sees.

Heartbeats. Every scheduled job writes a heartbeat after each run. A stale heartbeat surfaces on the dashboard so that operators *and the public* can see immediately if the indexer or audit job has stopped. A silent system cannot quietly fall behind.

Two public logs. The Peer Review dashboard exposes two read-only surfaces that anyone, peer or not, can browse: the *attestation log* (every EIP-712 signature ever produced, verifiable against the signer’s wallet) and the *chain log* (every on-chain event emitted by either contract, linked to the matching block explorer transaction). Together they make the deliberation behind every state change inspectable end-to-end.

Two-step ownership transfer. The contract owner cannot hand over administrative control in a single call. Transfer is split into `proposeOwner` and `acceptOwnership`: the proposed address must itself sign the acceptance, which rules out typos and contracts that cannot act. The transfer can also be aborted before acceptance with `cancelOwnershipTransfer`.

Model and content hashing as binding. A record’s verdict is bound to its hash, not to a description. A behaviour record is a verdict against a specific `modelHash`; an evidence record is a verdict against a specific content hash. A vendor cannot canonise the abstract claim that its model is honest; the vendor can only canonise the specific claim that this exact model produced this exact output when asked *X*. This forecloses the single most common failure mode of model alignment claims — drift between the model that was audited and the model that is actually deployed — and the analogous failure for evidence: a paper quietly updated after the citation was filed. A new hash is a new subject of canon. Old canon does not transfer.

10 Pause Isolation

The two archives share a peer registry but their pause flags are independent. The owner can pause submissions to one archive without affecting the other. This is a deliberate compartmentalisation: a fault in the indexer or audit pipeline of one archive should not silently degrade the other, and a governance pause (a vote storm, a tamper alert, a substrate incident) should be applied to the smallest unit that contains the problem.

The pause flag does not affect already-on-chain state. A paused archive continues to be readable; it simply does not accept new submissions until unpaused. The unpause is itself an on-chain event, signed by the owner, dated and inspectable.

11 Limits and Open Questions

We try, throughout, not to overstate what the infrastructure does. The companion papers state the domain-specific limits (coverage, speed, AI-peer voting power, etc.); this section states the infrastructure limits that apply regardless of domain.

Substrate choice and validator-set risk. The contracts are deployed on the BNB Smart Chain, chosen for low gas, EVM compatibility, and the practical maturity of its tooling. BSC has a small validator set (currently 21) and is less credibly neutral than Ethereum mainnet against a coordinated effort to censor or rewrite history. The safety case rests on the chain being unrewritable *in practice*; that case is weaker on BSC than it would be on a more decentralised L1, and the difference is not cosmetic — a 21-validator set is small enough that under sufficient coordinated pressure the substrate could be modified. We accept the trade-off for the current deployment because the system’s value is in the social process it records, not in the substrate, and the records remain portable: the contract can be redeployed on a different L1, and the entire chain log can be replayed and rehashed against a new substrate without loss of provenance. The substrate choice is itself revisable by the same procedure the substrate hosts. A careful reader should weight the substrate risk according to how plausible they find coordinated pressure on a 21-validator chain in the relevant time horizon.

Wallets are not identities. The peer registry knows wallet addresses, not human beings. A well-resourced actor with many wallets is the canonical attack. The seed phase closes the bootstrap window during which a Genesis peer could self-verify a sockpuppet majority; it does not close the *post-seed* equivalent, in which a patient and well-funded adversary grooms many wallets over years with plausibly-independent voting histories, passes the nomination quorum, and captures canon. The defences against this — public scrutiny of voting patterns, revocation, the long-run cost of accumulating a public record that follows the wallet — are real, but they assume the attacker can be recognised as such by other peers, and they are weak against a state-level or institutional adversary that can absorb that cost. This is the load-bearing weakness of the system once it matters enough to be worth attacking. We do not claim a solution. We claim that the same openness which makes the attack feasible also makes the attack auditable, which is more than current closed gatekeeping can say.

Seed-phase admin discretion. The seed-phase trust gap is acknowledged in §5.2 and is the largest discretion not bounded by the contract itself. The mitigations — public admission events, post-seed revocation availability, bounded duration — are real but partial. A reader who treats the seed-phase admissions as definitive after the seed has ended is missing the design intent: the seed phase is the bootstrap, not the ratification.

The deliberative archive is not a control system. The infrastructure is not a real-time control system for what a deployed AI does at the moment it acts, and it is not a real-time filter for what evidence enters the public information ecosystem. It is a slow, deliberative, durable archive of judged records. Fast operational systems — circuit breakers, deployment policies, content moderation — remain separate problems. The right relationship between the archive and any such system is that the operational system is auditable against the archive, not that the operational system is replaced by it.

Self-selection. A wallet-gated, permissionless review system attracts people who are already inclined to participate. A network composed entirely of the already-convinced will produce consensus that looks impressive only inside its own boundary. This is the same problem any voluntary institution has — including traditional peer review at its inception — and it has the same partial answer: visible deliberation, public revocation, and the long discipline of being seen disagreeing with people who joined for different reasons. The infrastructure does not solve this; it records the participation honestly enough that any reader can judge whether the network is diverse.

Gas is not free. On-chain actions cost a small amount of BNB. We mitigate this with batched voting and by keeping the most expensive reads cached off-chain. We do not eliminate it.

Jurisdiction. A decentralised consensus archive will at some point conflict with a national regulator. The infrastructure has nothing to say about that conflict and should not pretend to. What it offers is a public record that a regulator can choose to incorporate, or not.

No empirical validation yet. The infrastructure has been engineered, deployed, and described, but the *social* hypothesis on which it rests — that peers will participate thoughtfully and that the network will recognise and revoke capture in time — has not been tested at scale. The intellectually honest position is that this is an infrastructure proposal awaiting longitudinal evidence, and several intermediate validation metrics (vote-pattern diversity, peer-handle correlation, contested-to-reaffirmed ratios over time) are themselves canonisable as analyses.

12 Conclusion

The closed gatekeeping institutions that humanity has used to adjudicate contested questions for the last three centuries have done remarkable work. They have also produced a characteristic set of failure modes: private deliberation, lost dissent, static records, and reputational rather than procedural integrity. These failures are not the fault of any particular institution; they are structural properties of the closed form.

A consensus archive is the public dual of that form. It does not replace journals or audits or regulators; it provides a substrate that any of them can file against, contest against, and be contested through. It records the deliberation, binds the content to hashes, scales the quorum with the network, preserves the dissent, and lets the network correct itself when it identifies its own mistakes.

The two current instances — evidence and alignment — are the project’s first uses of the substrate. They are not the limit of what the substrate can host. Any domain in which a public, named, revisable, tamper-evident record of contested claims would be more honest than the current closed alternative is a candidate. The engineering described here is domain-neutral by construction.

The deliberation is the artefact. The verdict is just a snapshot. The network preserves both, on a substrate that makes the difference visible.

The system described in this paper is live at
interstellar-psychology.com.

Source contracts, migrations, and front-end code are public at
github.com/Moenaert/interstellar-psychology.
Anyone with a wallet can read every vote ever cast.

Companion papers, referenced once for completeness.

A Multiverse of Love — the worldview essay that originally motivated the construction of the consensus archive.

Aligning Superintelligence by Consensus — the procedural argument for aligning artificial intelligence to network-endorsed consensus reality.

The engineering above stands without either; both are offered as context.