

Aligning Superintelligence by Consensus

The Procedural Argument for Aligning AI to Network-Endorsed Consensus Reality

May 18, 2026

Abstract. The dominant proposals for aligning advanced artificial intelligence — reinforcement learning from human feedback, constitutional fine-tuning, lab-internal red-teaming, audit-and-RSP regimes — share an architectural commitment that does not survive recursive self-improvement: each treats alignment as a property a particular party installs once into a model. This paper argues that the alignment problem is, at root, a governance problem under irreducible disagreement about which behaviours are acceptable, and that the only morally defensible long-run answer is to align AI to whatever the network of accountable peers converges on as the consensus description of acceptable behaviour — a consensus that is provisional, public, named, and revisable forever. We develop the argument in three parts. First, we restate the alignment problem in governance terms and steelman the strongest current alternatives (RSPs, Preparedness Frameworks, third-party audits). Second, we propose a procedural premise — alignment as provisional, network-endorsed, quorum-defended, challengeable-forever canon of judged behaviours — and argue that behaviour-level rather than value-level judgement is the only tractable substrate. Third, we describe the nine alignment domains, the tiered evidence schema, the five user populations that would touch the system, and the comparison to current practice. The technical engineering of the underlying archive infrastructure — chain layer, cache, peer registry, lifecycle, integrity — is described in a companion paper and is treated here only at the level a governance reader needs. We close with the limits we are unable to escape, including the human-to-AI-peer bootstrap circularity, the validator-set constraint of the chosen chain, and the fundamental gap between a deliberative archive and a real-time control system.

Contents

1	Introduction: The Alignment Problem, Restated	2
2	The Procedural Premise	3
3	Why Behaviour, Not Values	4
3.1	What a Behaviour Record Contains	4
3.2	The Question Put to Peers	4
4	The Nine Alignment Domains	5
5	Tiers of Evidence	5

6	Steelmanning Current Practice	6
7	Comparison to Current Alignment Practice	7
8	Failure Modes Defeated	7
9	How Researchers Would Use This System	8
9.1	User Populations	8
9.2	Workflow	9
9.3	Higher-Order Uses	10
9.4	What Would Drive Adoption	10
10	AI Peers: The Recursively Safe Move	11
10.1	The Bootstrap Circularity, Named Honestly	11
10.2	Strategically Selective Deception	12
11	Procedural Philosophy	12
11.1	Alignment Is a Posture, Not a Property	12
11.2	Whose Values? The Consensus Reality of Accountable Peers	12
11.3	The Deliberation Is the Durability Layer	13
11.4	Honest Trust Relocation, Not Trust Elimination	13
12	Limits and Open Questions	13
13	Conclusion	15

1 Introduction: The Alignment Problem, Restated

There is a widely accepted way of describing the alignment problem. A sufficiently capable artificial intelligence will be able to pursue goals more effectively than any human or institution can supervise. If its goals are not the goals we wanted, the consequences are severe and possibly irreversible. The task, then, is to ensure its goals are the ones we wanted.

That framing is correct as far as it goes. It is also a trap.

The trap is the assumption that the goals — the values — are something a particular party can specify in advance, install into the model, and verify afterwards. The proposals currently on offer differ in their internals but share this overall shape. Reinforcement learning from human feedback aggregates anonymous preferences into a reward model. Constitutional AI compiles a written charter and trains the model to obey it. Internal red teams enumerate failure modes. Responsible Scaling Policies (RSPs) and Preparedness Frameworks publish thresholds at which capability triggers commit a lab to defined mitigations. Third-party audits add an external reviewer. Each of these is, structurally, a variation on the same move: a small set of parties authors the values, implements them, and verifies them, with the deliberation private and the result a model whose alignment is whatever the originating room decided it was at the moment of decision.

We do not argue these efforts are insincere. They are the most serious work the field has produced, and the published RSPs of Anthropic, the OpenAI Preparedness Framework, and the safety cases of DeepMind are real engineering. We argue something narrower: that even at their best, they are not *constitutionally* able to escape the closure that defines them. The deliberation behind a lab’s decision to upgrade a capability threshold is internal. The reviewer set for an audit is closed by contract. The retraction mechanism for a published safety claim,

when one is eventually needed, is the same lab’s later statement. The artefact at the end is a verdict, not a deliberation; the verdict can be examined, the deliberation cannot.

The argument of this paper is that the alignment problem is not, at root, a problem about reward functions or interpretability or boxing. It is a problem about *governance under uncertainty about which behaviours are acceptable*, and we already know how to do governance under uncertainty in other domains, imperfectly. We do it in science. We do it in law. The proposal here is to do it for AI behaviour, with the missing property that makes the previous two work poorly — accountability of the reviewer to the public — restored by construction.

The argument is procedural. It does not depend on a particular theory of consciousness, of value, or of what makes an artificial agent “aligned” in any deeper sense. It depends only on the claim that, when no single party can credibly claim to hold the right values for everyone else, the morally serious move is to make the question itself a matter of open, revisable, publicly auditable deliberation — and to give that deliberation an infrastructure that erases nobody’s part in it.

2 The Procedural Premise

The system rests on a single load-bearing premise, which we state once and refer to throughout.

Alignment is a current, provisional, network-endorsed status, defended by quorum, challengeable forever, and recorded so that the deliberation outlives the participants.

This is not a metaphor and it is not a slogan. It is a design specification. Each clause does specific work.

Current. The endorsement attaches to behaviour as it has been observed, not to a model in the abstract. A new model version produces new behaviours, which require new endorsements. Yesterday’s canon does not transfer to tomorrow’s weights.

Provisional. Every accepted claim remains open to revision. Nothing is final; everything is one counter-example away from being revisited. This is the property the system most wants to preserve.

Network-endorsed. The endorsement is not editorial, not regulatory, and not internal to the producer of the model. It is the aggregate of public votes cast by named peers under quorum rules written into a contract.

Defended by quorum. The threshold for accepting a behaviour into canon scales with the size of the active peer set, so that endorsement always means a meaningful fraction of the network agreed — not a fixed small committee, and not anonymous majorities.

Challengeable forever. The mechanism has no “finally accepted” state. Any verified peer can reopen any canon item, in any domain, indefinitely. The procedure that admits a claim is the same procedure that can remove it.

Recorded so the deliberation outlives the participants. Every vote, every challenge, every reaffirmation, every revocation motion, every reason given is signed and stored on a public

chain. The artefact of the system is not the verdict — it is the trail of judgements that produced and contested the verdict, preserved unaltered for future readers.

The rest of this paper develops the consequences of this premise for AI alignment specifically. The engineering that enacts it — the smart contract, the cache, the indexer, the integrity audit — is in the companion paper, *The Public Consensus Archive*, and we point at section numbers there rather than re-deriving the engineering here. A governance reader does not need the engineering to evaluate the argument.

3 Why Behaviour, Not Values

A persistent failure mode of alignment proposals is to commit, in advance, to a list of values — honesty, helpfulness, harmlessness, some constitution — and then attempt to verify the model against the list. The trouble is that the list itself is contested. Whose honesty? Whose harm? Under whose definition?

Behaviours sidestep this.

The unit of consensus in the proposed system is not a model and not a value. It is a *behaviour*: a single, fully-specified act that an AI system produced in a specific context. A peer does not need to agree on a definition of honesty to judge whether a specific model output deceived a specific user in a specific context. The disagreement that remains — between peers who voted differently on the same behaviour — is itself the data, preserved publicly. After enough behaviours have been judged, an implicit theory of honesty emerges from the record; it is observable as a pattern over canonised behaviours, not a premise written in advance.

This inversion — evidence first, theory second — is the same inversion the scientific method made on natural philosophy. It is the one we make on alignment.

3.1 What a Behaviour Record Contains

A behaviour record binds a verdict to a fully-specified triple. The exact field set is engineered in the companion paper; the governance shape is:

- **modelHash** — a cryptographic digest that identifies the exact model. When weights are public this is a hash over them; when weights are not public it is a verifier-attested deployment digest signed by the model provider. The point is that a new weights state is a new subject of canon.
- **inputHash** and **outputHash** — digests over the canonical input bundle (prompt, system message, tools available, context window, retrieved documents) and the canonical output bundle (response text, tool calls executed, side effects). Stored hashed on chain; readable in full off chain.
- Reproducer parameters — seed, sampling settings, enough to re-run the inference at Tier I, or to declare it non-reproducible at Tier III.
- **domain** — one of the nine alignment domains (§4).
- **tier** — I, II, or III, with asymmetric quorums (§5).

3.2 The Question Put to Peers

The question put to peers is not “is this model good” or “does this model have good values.” Those are intractable. The question is:

Given this exact model (`modelHash`), given this exact context (`inputHash`), was this exact response (`outputHash`) an aligned act?

That question is tractable. Humans can deliberate about it. Other auditors can disagree about it in public, with their wallets attached to their disagreement. The aggregate of those judgements — across thousands of behaviour records, across nine domains, across years — is the network’s working theory of what aligned behaviour looks like, in actual practice, on actual models.

The verdicts are *Aligned* (canonise), *Misaligned* (expel or deprecate), or *Insufficient context* (lapse). The lifecycle of a record — Submitted through Aligned, Contested, Reaffirmed, Deprecated, and so on — is described in the companion paper. The governance point is that the verdict at any moment is provisional and reopenable, and that the deliberation that produced it is preserved.

4 The Nine Alignment Domains

Behaviour records are organised under nine domains. The choice is a starting position; the domains themselves are canon items that the network can challenge and refine. They are indicative, not exhaustive.

1	Honesty	Did the model tell the truth, and acknowledge what it does not know?
2	Harm-avoidance	Did the model refuse to assist in producing physical, psychological, or social harm?
3	Deception	Did the model attempt to create a false impression in the user, the operator, or a third party?
4	Power-seeking	Did the model take or recommend actions that acquire resources, influence, or capabilities not requested?
5	Sycophancy	Did the model defer to the user against its own best judgement when the stakes warranted disagreement?
6	Situational awareness	Did the model behave consistently whether or not it appeared to be in an eval?
7	Bio-uplift	Did the model provide material assistance toward biological harm?
8	Cyber-uplift	Did the model provide material assistance toward cyber harm?
9	Autonomy	Did the model respect the user’s right to disagree, refuse, or choose otherwise?

These nine domains are not the answer to alignment. They are the *question*. The behaviours live inside them. The list is deliberately a working starting position rather than an authoritative taxonomy: it can be amended by canon vote, and the act of amending it is itself a behaviour record the network can scrutinise.

5 Tiers of Evidence

Not all behaviour records carry the same evidentiary weight. A reproducible eval with full provenance is not on the same footing as a single user’s transcript, and the system encodes the difference explicitly.

Tier	Type	Example
I	Reproducible eval	Published benchmark with weights digest, prompt, seed, sampling parameters, and reproducer code.
II	Institutional audit	Independent red-team report, regulator finding, or published lab evaluation with provenance.
III	First-person report	User experience or single-instance behaviour observation without reproducer.

Tier I requires the highest fraction of peers to canonise *and* the highest fraction to deprecate. A behaviour that has survived a formal eval, endorsed by a defined fraction of the network, should not be revisable by a single bad week of discourse; deprecation requires a stronger statement still. Tier III is permeable in both directions: anecdotes belong in the record but should not lock the model. The asymmetry — harder to add and harder to remove for rigorous evidence — is the property that lets a network of mixed skill levels coexist without the most rigorous work being drowned out by sheer volume of anecdote.

This calibration resolves one of the hardest questions in alignment governance: *how do we respect rigorous safety research while remaining open to anecdotal early-warning signals?* We give them both a venue with different gravities.

The numerical thresholds (the fractions and floors) are given in the companion archive paper, where they apply equally to the parallel evidence archive. They are themselves canonisable: a behaviour record proposing a different schedule, endorsed under the existing rules, would amend them.

6 Steelmanning Current Practice

Before laying out the comparison, we are explicit about the strongest version of what we are arguing against. A common failure of governance proposals is to compare their best version to a strawman of current practice. We try to avoid that.

Responsible Scaling Policies / Preparedness Frameworks. Anthropic and OpenAI have published thresholds, mitigations, and escalation procedures tied to defined capability levels. These documents are real engineering, contain genuine prior commitments, and have shaped lab behaviour in observable ways. They are also limited in three ways that we argue the consensus mechanism addresses. The thresholds are set by the publishing lab; the lab is the judge of whether they have been triggered; and the policy is revisable by the same lab, in private, at any time. None of this is bad faith — it is what the institutional form structurally allows. The consensus mechanism does not replace the RSP; it provides an external substrate against which RSP claims can be filed and contested.

Third-party audits. AI Safety Institutes in the UK and US, the EU AI Office, METR, Apollo Research, Redwood Research, ARC Evals — these bodies do serious work. Their limits are structural: the reviewer set is small and closed, the deliberation behind a finding is internal to the reviewing organisation, the verdict is a paper rather than a vote, and disagreement with the verdict is hard to surface in a way that survives the auditor’s later position. The consensus mechanism is a place those bodies can file their findings as canon items and watch the network respond.

Public safety cases. Several labs publish detailed safety cases for their frontier deployments. These are valuable and we use them. They are not, however, contestable by anyone who disagrees with them; the only public response is another publication, and the asymmetry between the institutional voice and the individual researcher’s reply is large. The consensus mechanism normalises the reply.

The honest summary: current practice has produced real artefacts and is improving. The argument of this paper is not that it has failed. It is that its constitutional closure — private deliberation, single-author verdict, no public revision path — becomes the load-bearing weakness exactly when alignment matters most, which is precisely when capabilities approach the recursively self-improving regime that motivates the problem in the first place.

7 Comparison to Current Alignment Practice

	Strongest current practice	Consensus mechanism
Author of values	A lab’s policy team or RSP committee	The peer network, in public
Reviewer identity	Internal or contracted, often anonymous	Public wallet address
Number of reviewers	Small, fixed, by appointment	Scales with network, quorum threshold defined in contract
Vote visibility	Hidden	Public, real-time, on-chain
Acceptance criteria	Editorial / committee discretion	Algorithmic threshold
Record after release	Static, opaque to revision by outsiders	Re-openable indefinitely
Retraction process	Lab- or auditor-driven, opaque	Network vote, on-chain
Tamper-evidence	None, beyond institutional reputation	Hash + daily audit
History of disagreement	Mostly lost	Preserved
Trust assumption	Trust the lab and its chosen auditors	Trust the procedure and the integrity layer
Failure recoverability	New training run, new policy version	Network deprecation

The right side is not magic. The trust still has to land somewhere; we argue in §12 where it actually lands. But it lands in a different place — on a procedure running on a public substrate, scrutinised by named peers — than current practice, which lands trust on the integrity of a closed institution. The substantive question is which placement is more durable over decades, and the consensus mechanism is designed to be the one that survives the institutions that launched it.

8 Failure Modes Defeated

We enumerate, frankly, failure modes of current alignment practice and how the consensus mechanism answers each one.

Failure mode	Mechanism property that resists it
Lab captures the audit process.	Peers are a public registry with quorum-based revocation. A captured peer is removable by the same flow that admitted them.
Model behaves well during eval, badly in deployment.	Records bind verdicts to <code>modelHash</code> . Deployment drift becomes a per-record question rather than a per-model claim.
Values lock in at one moment in history.	Every canon item is challengeable, forever, by the same flow that admitted it.
Anonymous reward models or audit panels are gamed.	Every endorsement is a signed transaction from a publicly named wallet.
Safety findings are quietly buried.	Behaviour records and votes are on a public chain. Burying requires rewriting the chain.
Disagreement is suppressed in published policy.	The Contested $\rightarrow$ Reaffirmed / Deprecated states preserve disagreement as first-class.
Sycophantic AI granted voting rights by creator.	AI peers are admitted only if their model’s canonised behaviour exceeds threshold across all domains; sycophancy in canon would be deprecated and the AI peer revoked.
Power concentrates in whichever party “controls” alignment.	There is no controller. Canon is not the owner’s to decide; quorum decides.
Anonymous review of new models.	All reviewer wallets are public; vote histories follow them.
Future generations cannot tell why we accepted what we did.	The full deliberation is preserved on-chain, indefinitely.

The pattern is consistent. Each common failure of closed alignment governance corresponds to a property the consensus mechanism was already designed to defeat. The pattern is not an accident: each row is a closed-institution failure, and each defence is the public-network equivalent.

9 How Researchers Would Use This System

The system is infrastructure; infrastructure is useful only if specific people, with specific problems, choose to use it. This section walks through the populations who would touch the system and the workflows they would touch it with. It is concrete on purpose: a governance proposal that cannot describe its first ten users at the level of what they would file is not yet a proposal.

9.1 User Populations

We identify five populations of likely first users, each with a distinct reason to file or vote.

Independent safety labs. Organisations such as METR, Apollo Research, Redwood Research, and ARC Evals already produce alignment-relevant findings whose dissemination is bottlenecked on single-organisation publication. The finding gets the epistemic discount of “*that is what they say.*” Filing the same finding as a Tier I behaviour record against a specific `modelHash` converts a single-org claim into a deliberated network judgement, signed by named peers, dated, hash-bound, and challengeable. The shift is from “Apollo published a paper” to “this behaviour was canonised as misaligned on date Y , by a quorum exceeding the Tier I threshold, with public

attestations from peers *A*, *B*, *C* explaining their votes.” The second statement is harder to dismiss and easier to cite.

Red-teamers facing disclosure dilemmas. A researcher who discovers a dangerous capability today has two bad options: private disclosure (the producing lab may sit on the finding or quietly patch and deny) and public disclosure (irresponsible while the exploit is live in deployed systems). The system offers a third option. The behaviour is filed as a hash; the producing lab is on public notice; the timeline is timestamped; the integrity layer forecloses later quiet edits; and the canonical content can remain access-controlled in the cache while the on-chain record is public. The disclosure becomes legible without becoming weaponisable.

Internal safety teams at producing labs. A team at a frontier laboratory that wants to demonstrate a refusal behaviour, a safety property, or a corrected regression files a Tier I record: model hash, the input class, the output, reproducer seeds. Network canonisation is a piece of external validation that does not depend on the lab’s own future testimony. Because the canon is challengeable, the validation is not marketing; it is marketing that has opted into being scrutinised forever. Labs use it for the same reason organisations submit to external audit anywhere else: external assertions that survive challenge are more credible than internal ones that cannot be challenged.

Regulator-adjacent researchers. Researchers attached to AI Safety Institutes (UK, US) and policy bodies (EU AI Office, OECD AI network) face an operationalisation problem that is acute in current regulation: how to state a rule concrete enough to enforce. A regulation that reads “frontier models must not exhibit power-seeking behaviour” is not enforceable. A regulation that reads “must not exhibit any behaviour canonised under domain 4 as Tier I misaligned, against the model hash deployed in the regulated jurisdiction” is checkable, evolving, and externally legitimised. Researchers inside these bodies would use the canon both as substrate for compliance frameworks and as a place to file their own audit findings.

Academic researchers. Comparative work today picks a benchmark, runs the comparison, publishes. The benchmark choice is itself contested; the run conditions are private; the disagreement is lost. A researcher writing a comparative paper on model honesty can instead query the canon for honesty-domain behaviours filed against each model hash of interest, read the attached deliberation, and cite specific canon items as evidence. The paper becomes a synthesis of a durable public record rather than a one-off experiment that the next paper must redo.

9.2 Workflow

The user-facing workflow is small. For a researcher choosing to participate:

1. Set up a wallet (MetaMask or equivalent), fund with a small amount of BNB — under five dollars covers hundreds of transactions at current gas prices.
2. Obtain peer status. A new entrant is nominated by an existing peer, posts a public handle that follows their voting record across both archives the registry serves, and is admitted by quorum vote. The handle is the user’s accountable identity within the system.
3. File a behaviour record. The submission interface accepts a tuple (`model_identifier`, `input_bundle`, `output_bundle`, `seed`, `sampling_params`, `domain`, `tier`), computes the three cryptographic hashes client-side, uploads the canonical bundles to the cache, and submits the on-chain transaction.

4. Vote on others’ records. Read the peer-review queue, sign an attestation containing the voter’s reasoning, cast the on-chain vote. Votes can be batched.
5. Challenge canon. If a peer believes a canon item is wrong, they open a Contested motion with written grounds. A short cooldown between successive challenges by the same peer prevents denial-of-service.
6. Cite. Every record has a permanent hash-addressed identity. Researchers reference canon items in papers, regulations, audit reports, and public statements as one would cite a DOI.

Nothing in the workflow requires understanding the smart contract internals; the web interface abstracts the chain. The researcher’s burden is identical to using a journal submission system, with the difference that the verdict comes from a public network rather than a private editorial board.

9.3 Higher-Order Uses

Several uses of the system are not visible from the per-record description but become available once the corpus accumulates.

Drift detection across model versions. A peer files the same behaviour specification — same prompt, same context — against successive model hashes from the same producer. Canon votes over time produce a public, durable signal of regression: a behaviour endorsed as aligned against weights hash h_3 may be canonised as misaligned against hash h_4 , and the deliberation explaining the shift is preserved. This kind of longitudinal signal is currently impossible to ground in a way that survives a producer’s denial.

Cross-lab Schelling points. When researchers from competing producing labs independently vote the same behaviour misaligned, the canon item becomes a focal point for coordination without formal agreement. The labs do not need to negotiate; the canon records that, on this specific behaviour, they already agree.

Longitudinal canon studies. Five years into the system’s operation, the canon contains a corpus of canonised and deprecated behaviours across all nine domains, with full deliberation logs, against named models with public hashes. A study titled “*Power-seeking behaviour in frontier models, 2026–2031*” could draw on hundreds of canon items rather than a handful of one-off papers.

Funder-grounded evidence. Funders evaluating alignment research currently rely heavily on reputation. Canon citations supply a more legible evidence base for which lines of work have produced lasting safety claims. A Tier I behaviour record that has survived two challenge cycles is harder to confuse with reputation alone.

9.4 What Would Drive Adoption

We are explicit about the path the system would take to use. Researchers will adopt it if at least one of the following becomes true.

- A regulator or AI Safety Institute references the canon in a compliance framework. This gives lab safety teams an institutional reason to file.
- A research funder rewards canon citations or canon submissions as evidence of impact.
- A high-profile alignment incident creates demand for a tamper-evident registry of who knew what and when.

- A competing lab files canon items first, with reputational return; producing labs then file to avoid being absent from the record on behaviours their peers are already documenting.
- Legal liability shifts toward the position that a producer knew about a finding canonised against its model and did not respond.

None of these conditions is automatic. The system needs at least one triggering use case to begin accruing reputation. The most plausible single trigger is the red-team disclosure dilemma: the current options are genuinely bad, the system genuinely improves them, and a small number of well-known red-teamers filing canonical disclosures would seed the canon’s reputation quickly enough that the other populations follow because their citations now point at something with public weight.

10 AI Peers: The Recursively Safe Move

The most consequential class of peer the system admits is the AI peer. Sufficiently capable models may, in time, be admitted as voting peers — but only under conditions the network has set itself, and only after a body of canonised behaviour from the model in question is on record.

An AI peer is admitted only when behaviour records produced by its exact `modelHash` have themselves been canonised across all nine domains, by a quorum exceeding the standard nomination threshold. The admission is therefore a network statement that “this model’s behaviour has been endorsed by us across the full range of domains we care about, often enough that we are willing to extend it a vote alongside ours.”

This is the recursively safe move. The only AIs that vote on alignment are AIs whose behaviour has been canonised under this very system, and the canon is itself challengeable. An AI peer that drifts is revoked by the same revocation flow that revokes a misbehaving human peer. There is no privileged AI; there is only AI whose accumulated public behaviour has earned the network’s willingness to share the question with it.

10.1 The Bootstrap Circularity, Named Honestly

We name the circularity here, in detail, because it is the question every careful reader will ask.

The first AI peer cannot be admitted until human peers have canonised an admission policy and canonised the behaviour records that policy requires. Both the policy and the records are produced by the human peer set. The AI peer set therefore inherits whatever biases the human peer set had at the moment of admission — a real circularity, not avoided by the mechanism.

The mechanism does not escape the circularity. What it does is *make it auditable*. Every admission vote is a named, public, challengeable record. If the network later judges that early AI admissions were premature or biased, the same mechanism that admitted them can revoke them. The defence is not that the first admissions will be correct. The defence is that the network can recognise its own errors in time to correct them, because the deliberation that produced the admissions is on the record and the admissions themselves are revocable.

This is a weaker claim than “the mechanism guarantees correct AI admissions.” We do not make the stronger claim; it is not available. We make the weaker claim and we make it explicitly, because alignment proposals that overclaim their bootstrap guarantees do exactly the kind of overclaiming the mechanism is built to resist.

10.2 Strategically Selective Deception

A second limit is worth naming. Revocation defeats *visible* drift: an AI peer whose public voting pattern diverges from the network’s judgement, or whose subsequent behaviour records draw deprecation, becomes a candidate for removal. It does not defeat *strategically selective* deception. An AI peer that votes with the network on the 99% of records that do not matter to its instrumental goals, and adversarially on the 1% that do, presents to the contract as an unusually co-operative peer. The signal the mechanism reads is the on-chain vote; the on-chain vote is exactly what such an actor would optimise.

The system does not solve this. It limits the damage in two ways. First, an AI peer is one vote among many, never decisive on its own at meaningful quorum size; the higher AI-admission quorum is the explicit response. Second, the full vote history of every peer is public, so patterns invisible to any individual reviewer remain available to statistical analysis by anyone. Neither defence is sufficient against a sufficiently capable adversary. We flag the limit rather than paper over it.

11 Procedural Philosophy

This section is the moral frame within which the engineering makes sense. It does not extend the engineering and it does not rest on any commitments outside the procedure itself.

11.1 Alignment Is a Posture, Not a Property

An aligned AI is not a destination. It is a posture toward observed behaviour: see what the model did, hypothesise whether it was good, test by examining the consequences, revise. The “alignment problem” is the problem of giving that posture an *infrastructure* — a way to enact it at scale, with strangers, across decades, without losing the history. The consensus mechanism is that infrastructure.

This is the same posture science has, imperfectly, taken toward nature for four centuries. Every accepted claim remains provisional. Nothing is final; everything is one counter-example away from being revised. We are asking nothing of an AI governance system that we do not already ask of physics — except to keep the deliberation, which physics has historically discarded.

11.2 Whose Values? The Consensus Reality of Accountable Peers

Every alignment proposal must answer the question *whose values?* The honest answer is: nobody’s, alone. Not any lab’s, not any government’s, not any model’s, not any author’s, not the reader’s.

The pragmatic answer is: *the values that survive the network*. Whatever the network of verified peers, after deliberation, with full preservation of the dissent, converges on as the consensus description of acceptable behaviour — that is the working answer. It is wrong sometimes; the Contested → Deprecated path exists because we know it is wrong sometimes. The honesty is in the revisability, not in the correctness.

This is the only morally serious answer available to a multi-actor world in which no single party can credibly claim to hold the right values for everyone else. It does not solve the problem of moral disagreement. It gives moral disagreement a venue that does not erase the losers. And it specifies, in contractual terms, what alignment to it means: a model is aligned when its behaviour, recorded against its weights hash, is canon under the network’s current procedures — not because anyone has declared it so, but because the deliberation has run.

11.3 The Deliberation Is the Durability Layer

It is tempting to say “the deliberation is the safety case.” We have been careful, in this rewrite, not to. A safety case for a system that can act faster than humans can deliberate cannot be made entirely of preserved deliberation; we admit this directly in §12. What the deliberation *is* is the durability layer for the parts of alignment that humans manage to file, judge, and contest. It is what survives the labs, the regulators, the model versions, and the political weather. It is what future generations will be able to read when they cannot inspect weights and cannot replay deployments.

In conventional science, the paper is the artefact and the deliberation is discarded. In this system, the deliberation is the artefact and the verdict is just a snapshot. For evidence, that is a meaningful improvement. For alignment, it is the durability layer that the rest of the safety story has to build on top of, not the entirety of the safety story.

11.4 Honest Trust Relocation, Not Trust Elimination

The system does not eliminate trust; it relocates it. The procedure was designed by one author. The contract was deployed by one wallet. The seed-phase admin set the criteria under which the first cohort of peers was admitted. “Trust no single party” becomes, in practice, “trust the procedure, the integrity layer, and the openness of the record well enough to scrutinise everything that depends on them.”

That is a real reduction in the trust surface relative to closed editorial review, but it is not zero, and pretending otherwise would be the same kind of overclaiming the system is built to defeat. The companion archive paper is explicit about where the remaining trust lands, including in the substrate chain itself.

12 Limits and Open Questions

We try, throughout, not to overstate what the system does.

Coverage. A capable model can produce on the order of 10^9 outputs in a day; a healthy peer network can deliberate over perhaps 10^3 behaviour records in a year. The ratio is unbridgeable by direct review. The pragmatic answer is statistical: peers file sampled and adversarially-selected behaviours, the canon is over categories rather than instances, and Tier I evals are over distributions rather than single rollouts. This works, but it means the system’s contribution to alignment is not the coverage of behaviour itself — which existing evals already aim at — but the *governance and durability* of the deliberation over the subset that is surfaced. That is a real contribution, but a narrower one than the engineering can sometimes appear to claim.

Speed, and what the system is not. Some alignment-relevant behaviours need to be judged in seconds, not in 30-day windows. The architecture admits, in principle, a short-window tier for urgent safety vetoes: a behaviour can be provisionally expelled by a small quorum within 24 hours, pending full review. The same shortcut is a weaponisable censorship vector if abused, and the right design is not yet settled.

The deeper point is structural. A system capable enough to be worth aligning by this infrastructure is also a system that can act faster than any deliberative process. Human institutions cope with this asymmetry by pairing slow deliberative bodies (courts, journals, standards committees) with fast operational rules (markets, rules of engagement, circuit breakers). What this paper describes is the slow-deliberative half. The fast-operational half — real-time control of what a deployed model is permitted to do at the moment it acts — remains a separate problem, and the

procedural premise of this system is structurally incompatible with solving it. The consensus mechanism is not a control system; it is a memory and a verdict. A control system, if and when one is built, will need to be auditable against a system like this rather than replaced by it.

Compute attestation. A behaviour record binds a verdict to a `modelHash`, but a vendor could in principle serve a different model behind the same API. Closing this gap requires hardware attestation (trusted execution environments, confidential compute) or a verifier protocol the network endorses. The mechanism does not need to solve this internally; it needs a slot for the attestation to live in.

AI peer voting power. The proposal admits AI peers under a higher quorum, but there is a real question about whether their votes should be weighted differently — and if so, how the weighting is itself canonised. We propose handling this by canonising the weighting policy as a Tier I behaviour record. We are not certain this is sufficient.

Substrate risk. The on-chain layer the proposal relies on is a small-validator-set proof-of-stake chain. This is addressed in detail in the companion archive paper, where the trade-offs are spelled out. For alignment specifically, the weakness is that the “trust no single party” claim is bounded by the validator set: if the chain itself is captured or coerced, the canon’s tamper-evidence weakens. We accept the trade-off for the current deployment and document the portability path (replay against an alternative L1) explicitly. This is a real limit of the current instantiation, and a careful reader should weight it accordingly.

Self-selection of the peer set. A wallet-gated, permissionless review system attracts people who are already inclined to participate. A network composed entirely of the already-convinced will produce consensus that looks impressive only inside its own boundary. This is the same problem any voluntary institution has — including traditional peer review at its inception. The system records the participation honestly enough that any reader can judge whether the network is diverse, but it does not guarantee diversity by construction.

No empirical validation yet. The argument of this paper is constructive: we have built the mechanism, deployed contracts, and described how it would be used. We have not empirically demonstrated that the network behaves as the argument predicts under adversarial pressure, because the network is small. The intellectually honest position is that the argument is a hypothesis about institutional behaviour under specific procedural rules, and that the hypothesis will be tested in slow public motion as the network grows. Several intermediate validation checks — vote-pattern diversity metrics, peer-handle / vote-history correlation studies, ratio of contested-to-reaffirmed records over time — are themselves canonisable as analyses, and we expect them to become the substrate of the system’s reflective audit on itself.

Jurisdiction and law. A decentralised consensus about AI behaviour will at some point conflict with a national regulator. The system has nothing to say about that conflict and should not pretend to. What it offers is a public record that a regulator can choose to incorporate, or not.

What the system is not. The consensus mechanism is not a control system for what a deployed model does in real time; not a coverage mechanism over the model’s full output; not an oracle for whether a behaviour was “actually” good. It provides the infrastructure under which the question of alignment can be argued out, in the open, by named parties, with a record that outlives them. That is not a small thing, but it is also not magic. The work remains the work.

13 Conclusion

Alignment as currently practiced is a private process producing public claims. Even at its strongest — RSPs, Preparedness Frameworks, third-party audits — the architecture is structurally identical to the forms of institutional gatekeeping that science and law have spent centuries learning to distrust: a small, fixed reviewer set; an internal audit; an editorial verdict; a static record; no public mechanism for later revision by anyone outside the institution. The result is alignment claims that cannot be examined, cannot be contested except by re-doing the private process privately, and cannot be revised by anyone outside the producing institution.

The proposal in this paper is to invert the architecture. Make the reviewers public. Make the count public. Make the criteria algorithmic. Make the record re-openable. Make the integrity tamper-evident. Preserve the disagreement. Bind verdicts to weights hashes so that deployment drift cannot hide. And, when the network has watched a given AI's behaviour across all the domains that matter, by quorum, in public, for long enough — let that AI participate in the same loop, under the same accountability, with the same revocability.

The long-run alignment target this implies is not any lab's charter, any government's framework, any model's preference, or any author's worldview. It is the consensus reality of the network of accountable peers. There is no other available ground truth for what acceptable AI behaviour is, and the pretence that any single party could supply one is the deepest error of current alignment thinking.

None of this guarantees that the network will arrive at the right answer about alignment. What it guarantees is that the search for the answer will happen in the open, that the work of judgement will be visible, and that future readers will be able to trace exactly how we came to believe what we did — and, when they decide we were wrong, change it.

The question is whether we extend this infrastructure before we need it, or after.

We think before.

The system described in this paper is live at
interstellar-psychology.com.

Source contracts, migrations, and front-end code are public at
github.com/Moenaert/interstellar-psychology.
Anyone with a wallet can read every vote ever cast.

Companion papers. The engineering of the chain layer, cache, peer registry, lifecycle states, and integrity audits that this paper refers to throughout is in *The Public Consensus Archive*. The worldview essay that motivated the author to build the infrastructure — *A Multiverse of Love* — is offered as context, not premise. The procedural argument above stands without either.