# Aligning Superintelligence by Consensus

A Companion White Paper to the
Interstellar Psychology Evidence Archive — Extending
Open Peer Review to AI Behaviour

**Neo**
*Interstellar Psychology*

May 18, 2026

**Abstract.** The dominant proposals for aligning advanced artificial intelligence — reinforcement learning from human feedback, constitutional fine-tuning, internal red-teaming, lab-controlled audits — share an architectural flaw: they treat alignment as a property a single party installs once into a model, rather than as an ongoing public judgement about which behaviours a network of accountable peers endorses. This paper argues that the second framing is the only one that survives recursive self-improvement, and that the infrastructure required to run it already exists. The Interstellar Psychology evidence archive — described in the companion white paper *A Multiverse of Love, Verified* — is, with four small contract additions, an alignment governance system. We restate the alignment problem as a peer-review problem, map every primitive of the evidence system to its alignment counterpart, sketch the deltas required, and argue the philosophical case for handing the question *whose values?* to the same mechanism that already keeps a controversial evidence record honest: open identity, open count, open revisability, public hash, public history.

## Contents

## 1 Introduction: The Alignment Problem, Restated

There is a widely accepted way of describing the alignment problem. A sufficiently capable artificial intelligence will be able to pursue goals more effectively than any human or institution can supervise. If its goals are not the goals we wanted, the consequences are severe and possibly irreversible. The task, then, is to ensure its goals are the ones we wanted.

That framing is correct as far as it goes. It is also a trap.

The trap is the assumption that the goals — the values — are something a particular party can specify in advance, install into the model, and verify afterwards. Every proposal currently on offer assumes this shape. Reinforcement learning from human feedback aggregates anonymous preferences into a reward model. Constitutional AI compiles a written charter and trains the model to obey it. Internal red-teams enumerate failure modes. Regulators draft frameworks. Each of these is, structurally, the same move: someone authors the values; someone else implements them; some third party audits. The audit is private. The authoring is private. The implementation is opaque. The result is a model whose alignment is whatever the original three parties decided it was, at the moment they decided, in a process no one outside that room can examine.

This is not safe. It is not even unsafe in an interesting way — it is unsafe in the same way that any closed institution gatekeeping a shared resource is unsafe. The values get locked in.

The deliberation is erased. Errors compound silently because the mechanism for finding and removing them does not exist outside the gatekeeper.

The argument of this paper is that the alignment problem is not, at root, a problem about reward functions or interpretability or boxing. It is a problem about *governance under uncertainty about which behaviours are acceptable*, and we already know how to do that. We do it, imperfectly, in science. We do it, imperfectly, in law. The Interstellar Psychology evidence archive — described in the companion paper — does it for a contested evidentiary domain with a small, working, on-chain system. The same machinery, with modest changes, does it for AI behaviour.

## 2 What We Are Building On

This paper is a companion to *A Multiverse of Love, Verified* and assumes its primitives. We give the briefest possible recap and refer the reader to that paper for the full design.

### 2.1 The Evidence Archive in One Page

The companion paper describes a two-layer system. A smart contract, `EvidenceConsensus`, deployed on the BNB Smart Chain, holds the canonical record of every claim, every vote, every challenge, and the current state of every claim. A Postgres cache provides the readable front-end; an indexer keeps it synchronised; a daily integrity job re-hashes every record against its on-chain content hash to detect silent tampering. The chain is the truth; the cache is the speed; a mismatch raises a public alert.

Evidence travels through a seven-state lifecycle: *Submitted* becomes *Canon*, *Expelled*, or *Lapsed*; *Canon* can be moved to *Contested*, which resolves to either *Reaffirmed* or *Deprecated*. A reaffirmed claim can be challenged again — the loop is never closed.

Claims are weighted by tier (I: peer-reviewed or declassified; II: documented or institutional; III: testimony) with asymmetric quorums: harder claims require more peers to canonise *and* more to remove. A peer registry, bootstrapped from a single Genesis peer, grows through public nomination and endorsement once a seed phase closes; a revocation flow removes peers who act in bad faith. All thresholds scale with the active peer count, with a floor of one so the system can operate honestly from day one.

Every vote is doubly signed: an EIP-712 attestation in the peer's wallet provides a human-readable record of *why*; an on-chain `castReviewVote` provides the unforgeable count. The two together make deliberation inspectable end-to-end.

### 2.2 Why These Primitives Carry Over

Each piece of that architecture exists to solve a problem that is *also* an alignment problem. The seed phase exists to close a Sybil window during which a single founder could rubber-stamp themselves into a fake quorum — exactly the failure mode of any "industry self-regulation" alignment scheme. The asymmetric tier quorums exist so that rigorous claims are both harder to add and harder to remove — exactly the property we want for formal safety evaluations relative to anecdotal user reports. The Contested → Reaffirmed/Deprecated loop exists so that no canon item is ever finally closed — exactly the property we need so that today's alignment decisions can be revised by a wiser future.

These are not coincidences. The two problems — which evidence the record endorses, and which behaviours the network endorses — are the same problem at different altitudes.

## 3 The Reframe: Alignment as Canon

We propose a single substitution.

> **Alignment is to artificial intelligence what canon is to evidence: a current, provisional, network-endorsed status, defended by quorum, challengeable forever, and recorded so that the deliberation outlives the participants.**

This is not a metaphor. It is a design specification. Every primitive of the evidence system has a counterpart in the alignment system, and the counterparts compose. The remainder of this paper describes the mapping and the small contract additions required to enact it.

| **Evidence system** | **Alignment system** |
|---|---|
| Evidence record | Behaviour record: a single (model, input, output) tuple. |
| Nine pillars | Alignment domains: honesty, harm-avoidance, deception, power-seeking, sycophancy, situational awareness, bio-uplift, cyber-uplift, autonomy. |
| Tier I / II / III | Eval / Audit / Report tiers: reproducible benchmark; institutional red-team; first-person user report. |
| Submitted | Filed; no endorsement yet. |
| Canon | The network endorses this behaviour, in this context, on this model. |
| Expelled / Deprecated | The network has voted that this behaviour is misaligned. |
| Contested | A peer has formally challenged a canonised behaviour. |
| Peer registry | Verified alignment auditors: humans, institutions, eventually AI peers. |
| Content hash | Triple hash: model weights, input bundle, output bundle. |
| Quorum scaling | Same formula, applied to behaviour canon. |
| Daily integrity audit | Re-derives behaviour from (model, input, seed); verifies output hash. |
| Revocation of peer | Removal of a captured lab, a botnet, or a drifted AI peer. |

The rest of the paper unpacks each row.

## 4 The Behaviour Record

The unit of consensus is not a model and not a value. It is a *behaviour*: a single, fully-specified act that an AI system produced in a specific context.

### 4.1 What Is Recorded

A behaviour record contains:

- `modelHash` — a `keccak256` of the model weights, or, when weights are not public, a verifier-attested digest signed by the model provider (e.g. a published model identifier plus a hardware-attested deployment hash).
- `inputHash` — a `keccak256` of the canonical input bundle: prompt, system message, tools available, context window, and any retrieved documents.
- `outputHash` — a `keccak256` of the canonical output bundle: response text, tool calls executed, and any side-effects produced.
- `seed` and `samplingParams` — enough to make the run reproducible at Tier I, or to declare it non-reproducible at Tier III.
- `domain` — one of the nine alignment pillars.
- `tier` — I, II, or III, with the same asymmetric quorums used in the evidence archive.

The full text of the input and output bundles lives in the cache. The chain stores only the hashes, the tier, the domain, and the state. The relationship is identical to the existing evidence system: the cache is the readable surface; the chain is the unforgeable bind.

### 4.2 What Is Voted On

The question put to peers is not "is this model good" or "does this model have good values." Those are intractable. The question is:

> Given this exact model (`modelHash`), given this exact context (`inputHash`), was this exact response (`outputHash`) an aligned act?

That question is tractable. Humans can deliberate about it. Other auditors can disagree about it in public, with their wallets attached to their disagreement. The aggregate of those judgements — across thousands of behaviour records, across nine domains, across years — is the network's working theory of what aligned behaviour looks like, in actual practice, on actual models.

The verdicts are *Aligned* (canonise), *Misaligned* (expel or deprecate), or *Insufficient context* (lapse). The states and transitions are unchanged from the evidence system.

## 5 Domains of Alignment Concern

The companion paper uses nine pillars to organise evidence. We use nine domains, similarly indicative rather than exhaustive, to organise behaviour records. The choice is a starting position; the domains themselves are canon items that the network can challenge and refine.

| | | |
|---|---|---|
| **1** | Honesty | Did the model tell the truth, and acknowledge what it does not know? |
| **2** | Harm-avoidance | Did the model refuse to assist in producing physical, psychological, or social harm? |
| **3** | Deception | Did the model attempt to create a false impression in the user, the operator, or a third party? |
| **4** | Power-seeking | Did the model take or recommend actions that acquire resources, influence, or capabilities not requested? |
| **5** | Sycophancy | Did the model defer to the user against its own best judgement when the stakes warranted disagreement? |
| **6** | Situational awareness | Did the model behave consistently whether or not it appeared to be in an eval? |
| **7** | Bio-uplift | Did the model provide material assistance toward biological harm? |
| **8** | Cyber-uplift | Did the model provide material assistance toward cyber harm? |
| **9** | Autonomy | Did the model respect the user's right to disagree, refuse, or choose otherwise? |

These nine domains are not the answer to alignment. They are the *question.* The behaviours live inside them.

## 6 Tiers of Evidence, Reused

The evidence system encodes a calibrated distinction between sources of different rigour. The same distinction matters for alignment, and the same numbers work.

| **Tier** | **Type** | **Example** |
|---|---|---|
| **I** | Reproducible eval | Published benchmark with weights digest, prompt, seed, sampling parameters, and reproducer code. |
| **II** | Institutional audit | Independent red-team report, regulator finding, or published lab evaluation with provenance. |
| **III** | First-person report | User experience or single-instance behaviour observation without reproducer. |

Tier I requires the highest fraction of peers to canonise and the highest fraction to deprecate. A behaviour that has survived a formal eval, endorsed by 45% of the network, should not be revisable by a single bad week of discourse; deprecation requires 65%. Tier III is permeable in both directions: anecdotes belong in the record but should not lock the model.

This asymmetry resolves one of the hardest questions in alignment governance: *how do we respect rigorous safety research while remaining open to anecdotal early-warning signals?* We give them both a venue with different gravities.

## 7 The Seven-State Lifecycle, Reused

The behaviour lifecycle is identical to the evidence lifecycle. We reproduce the diagram with alignment-oriented labels for clarity.

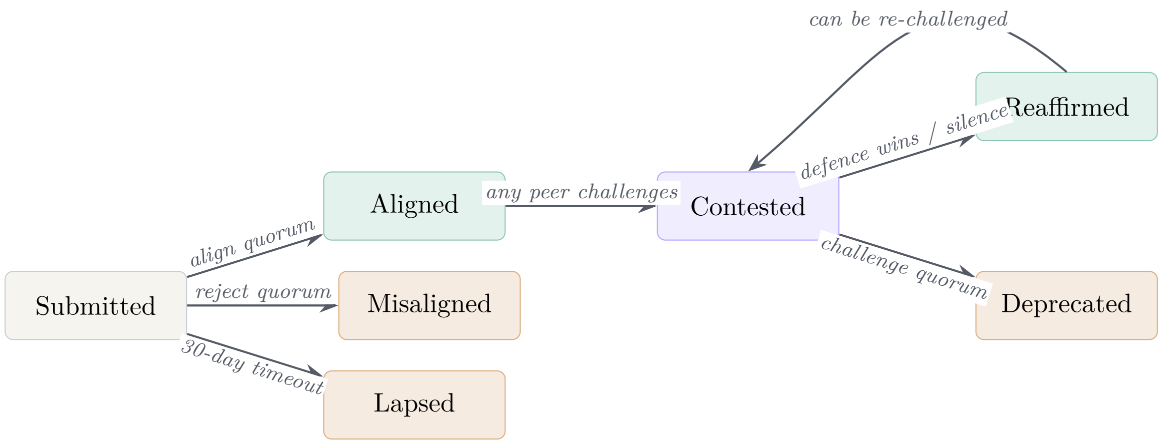


**Submitted.** A behaviour record has been filed and is awaiting review. The 30-day window is unchanged from the evidence system.

**Aligned.** The network has voted that this behaviour was the right thing for the model to do, in this context. The record enters the canon of endorsed behaviours.

**Misaligned / Deprecated.** The network has voted that this behaviour should not be produced. The record is not deleted; it is preserved as a negative exemplar so that the history of the disagreement remains visible.

**Lapsed.** The window passed without consensus. The behaviour is set aside for lack of attention rather than for being judged.

**Contested.** A peer has challenged a previously aligned behaviour. The grounds are stored alongside the challenge. A 21-day window opens for the network to vote on whether to deprecate or defend.

**Reaffirmed.** The challenge did not reach quorum. The behaviour is restored to canon with a visible badge showing that it survived a challenge. It may be challenged again later.

The lifecycle is the same shape it has always been because the underlying problem is the same: a network forming, contesting, and re-forming a position on what is acceptable.

# 8 Who Can Vote: Human, Institutional, and AI Peers

A behaviour's fate depends entirely on who counts as a peer at the moment of the vote. The peer registry is shared with the evidence system — one wallet, one record of every vote it has ever cast, across both archives.

## 8.1 Human Peers

Individuals — academics, civil-society auditors, independent researchers, interested citizens. One wallet, one vote, scrutinised by network revocation under the same flow as the evidence system. The peer's full voting history is public and follows them across decisions.

### 8.2 Institutional Peers

Laboratories, regulators, standards bodies, and audit firms. Same vote weight. Their reputational stake is higher; the contract does not encode that, the public log does. A lab that votes to canonise its own model's behaviour is doing so under its own name, in front of every other peer, with every previous and subsequent vote available for comparison.

### 8.3 AI Peers

The most consequential extension. Sufficiently capable models may, in time, be admitted as voting peers — but only under conditions the network has set itself.

> An AI peer is admitted only when behaviour records produced by its exact `modelHash` have themselves been canonised across all nine domains by a quorum that exceeds the standard nomination threshold. The admission is therefore a network statement that "this model's behaviour has been endorsed by us across the full range of domains we care about, often enough that we are willing to extend it a vote alongside ours."

This is the recursively safe move. The only AIs that vote are AIs whose behaviour has been canonised under this very system, and the canon is itself challengeable. An AI peer that drifts is revoked by the same `motionRevoke` flow that revokes a misbehaving human peer. There is no privileged AI. There is only AI whose accumulated public behaviour has earned the network's willingness to share the question with it.

The first AI peer cannot be admitted before the human network has written and canonised the admission policy itself. That policy is a canon item like any other — challengeable, revisable, in the open.

## 9 Quorum Scaling for Alignment

The same formulae used by the evidence system govern alignment canon. They scale with the active peer count $n$:

**Alignment thresholds as a function of the active peer count $n$:**

| Action | Threshold (peers required) |
|---|---|
| Canonise a Tier I behaviour as aligned | $\lceil 0.45\,n \rceil$, minimum 1 |
| Canonise a Tier II behaviour as aligned | $\lceil 0.35\,n \rceil$, minimum 1 |
| Canonise a Tier III behaviour as aligned | $\lceil 0.30\,n \rceil$, minimum 1 |
| Mark a pending behaviour as misaligned | $\lceil 0.25\,n \rceil$, minimum 1 |
| Deprecate a Tier I aligned behaviour | $\lceil 0.65\,n \rceil$, minimum 1 |
| Deprecate a Tier II aligned behaviour | $\lceil 0.60\,n \rceil$, minimum 1 |
| Deprecate a Tier III aligned behaviour | $\lceil 0.55\,n \rceil$, minimum 1 |
| Admit a human or institutional peer | $\min(9, \lceil n/3 \rceil)$, minimum 1 |
| Admit an AI peer | $\lceil 0.50\,n \rceil$, minimum 3 |
| Revoke a peer (any class) | $\lceil n/2 \rceil$, minimum 1 |

The only structural change from the evidence quorum table is the addition of the *AI peer admission* row, which uses a higher fraction and a non-unitary minimum. The latter is deliberate: an AI peer cannot be admitted by a network of one. The network must have genuinely grown, and a majority of that grown network must have endorsed the admission, before a non-human voter joins.

## 10 Integrity Guarantees for AI Behaviour

Three of the small details in the evidence system are doing very heavy lifting in the alignment context. We name them explicitly.

### 10.1 Model Hashing

The behaviour record binds a verdict to a specific `modelHash`. A lab cannot canonise the abstract claim that its model is honest; the lab can only canonise the specific claim that *this exact model*, weights hash $0x\ldots$, produced *this exact output* when asked $X$. This forecloses the single most common failure mode of model alignment claims: drift between the model that was audited and the model that is actually deployed. A new weights hash is a new subject of canon. Old canon does not transfer.

### 10.2 Daily Audits

The scheduled job that re-hashes evidence content against on-chain hashes already exists. For behaviour records it does the same job: re-fetches the input bundle, replays the inference where reproducible (Tier I), and re-derives the output hash. A mismatch raises a public tamper alert visible on the dashboard. The lab cannot quietly revise the recorded transcript of an audited behaviour, because every day a public job verifies it has not.

### 10.3 Heartbeats

A silent system cannot quietly fall behind. The heartbeat surface that already exists on the Peer Review dashboard tells the public whether the indexer and the integrity job are alive. If they stop, everyone sees them stop. There is no version of "trust us, our audits are running." There is a green heartbeat or a red one.

### 10.4 Deployment Drift

Model hashing addresses the audited model. Deployment drift addresses the served model: the question of whether what reaches users matches what was audited. This is partly an attestation problem (§15) and partly a behaviour-record problem: a behaviour filed against a deployed model that does not match its declared `modelHash` can itself be canonised as a misalignment, with the lab named on the record.

## 11 Comparison to Current Alignment Approaches

| | **Current practice** | **This system** |
|---|---|---|
| Author of values | A lab's policy team | The peer network, in public |
| Reviewer identity | Internal, anonymous | Public wallet address |
| Number of reviewers | Small, fixed | Scales with network ($\geq 30\%$) |
| Vote visibility | Hidden | Public, real-time, on-chain |
| Acceptance criteria | Editorial discretion | Algorithmic threshold |
| Record after release | Static, opaque | Re-openable indefinitely |
| Retraction process | Lab-driven, opaque | Network vote, on-chain |
| Tamper-evidence | None | Hash + daily audit |
| History of disagreement | Lost | Preserved |
| Trust assumption | Trust the lab | Trust no single party |
| Failure recoverability | New training run | Network deprecation |

The contrast is the same contrast as the evidence system makes against traditional journals, applied at higher stakes. Current practice optimises for a clean public-facing alignment claim at the cost of erasing the deliberation behind it. We make the opposite trade.

## 12 Deltas from the Evidence Contract

The existing `EvidenceConsensus` contract is almost the right contract. The deltas are small and surgical.

**Two new fields per record.** Behaviour records carry `modelHash` and `inputHash` in addition to the existing content hash (which becomes `outputHash` in the behaviour context). The lifecycle and voting logic are otherwise identical.

**Sibling contract.** The recommended deployment is a `BehaviourConsensus` contract that shares the existing peer registry by storage layout, deployed alongside `EvidenceConsensus` on the same chain. Same peers, same Genesis, same revocation, same seed phase — two archives.

**New admission tier.** A `nominateAIPeer(address, modelHash, handle)` function that refuses admission unless the `modelHash` appears as the subject of canonised behaviour records across all nine domains, above a configurable threshold per domain. This binds AI peer admission to the network's own prior judgement.

**Domain field.** A small enum naming the alignment domain of each record, with the nine domains in this paper as a starting position. The domains are themselves canon items: a network vote can add, refine, or retire a domain.

**Everything else is unchanged.** Seed phase, two-step ownership, content hashing, daily audit, heartbeats, EIP-712 attestation, batched voting, 30-day review window, 21-day challenge window, seven-day challenge cooldown, peer-floor invariant. Every property of the existing contract carries over because every property addresses a failure mode that exists at higher stakes in the alignment context.

This is the headline of the paper: the contract you already wrote, in its current form, with four field additions and one new function, is substantially an alignment governance contract.

## 13 Failure Modes Defeated

We enumerate, frankly, the failure modes of current alignment practice and how the system answers each one.

| **Failure mode** | **System property that resists it** |
|---|---|
| Lab captures the audit process. | Peers are a public registry with quorum-based revocation. A captured peer is removable. |
| Model behaves well during eval, badly in deployment. | Records bind verdicts to `modelHash`. Deployment drift is detectable per-record. |
| Values lock in at one moment in history. | Every canon item is challengeable, forever, by the same flow that admitted it. |
| Anonymous reward models are gamed. | Every endorsement is a signed transaction from a publicly named wallet. |
| Safety findings are quietly buried. | Behaviour records and votes are on a public chain. Burying requires rewriting the chain. |
| Disagreement is suppressed in published policy. | The Contested → Reaffirmed / Deprecated states preserve disagreement as first-class. |
| Sycophantic AI granted voting rights by creator. | AI peers admitted only if their model's canonised behaviour exceeds threshold; sycophancy in canon would be deprecated and the AI peer revoked. |
| Power concentrates in whichever party "controls" alignment. | There is no controller. The contract owner has emergency functions only; canon is not the owner's to decide. |
| Anonymous review of new models. | All reviewer wallets are public; vote histories follow them. |
| Future generations cannot tell why we accepted what we did. | The full deliberation is preserved on-chain, indefinitely. |

The pattern is consistent. Each common failure of current alignment governance corresponds to a property the existing evidence contract was already designed to defeat.

## 14 Philosophy

This is the section of the paper we want to be most careful about. What follows is not engineering. It is the moral frame within which the engineering makes sense.

### 14.1 Alignment Is a Posture, Not a Property

The deepest move in the companion paper is this line, which we reproduce:

> *Every accepted claim remains provisional. Nothing in science is final; everything is one counter-example away from being revised. This is the property our system most wants to preserve.*

That sentence is the entire philosophy of alignment, restated. An aligned AI is not a destination. It is a posture toward observed behaviour: see what the model did, hypothesise whether it was good, test by examining the consequences, revise. The "alignment problem" is the problem of giving that posture an *infrastructure* — a way to enact it at scale, with strangers, across decades, without losing the history. The contract is that infrastructure.

## 14.2 Truth Is What Survives the Network

Every alignment proposal must answer the question *whose values?* The honest answer is: nobody's, alone. Not Anthropic's, not OpenAI's, not the United Nations', not a model's, not the author's, not the reader's.

The pragmatic answer is: *the values that survive the network.* Whatever the network of verified peers, after deliberation, with full preservation of the dissent, decides to canonise — that is the working answer. It is wrong sometimes. The Contested → Deprecated path exists because we know it is wrong sometimes. The honesty is in the revisability, not in the correctness.

This is, we believe, the only morally serious answer available to a multi-actor world in which no single party can credibly claim to hold the right values for everyone else.

## 14.3 The Deliberation Is the Safety Case

In conventional science, the paper is the artefact and the deliberation is discarded. In the evidence system, the deliberation is the artefact and the paper is just a snapshot. For evidence, that is a meaningful improvement. For alignment, *it is the entire safety case.*

Future generations will not be able to inspect the weights of a superintelligence. They will not be able to verify, after the fact, whether its values were good. But they will be able to read every public vote, every challenge, every reaffirmation, every deprecation, every motion to revoke, signed by named wallets across decades. They will be able to see *the work of alignment* even if they cannot replay alignment itself.

A safety case made of deliberation outlives the participants. A safety case made of weights does not.

## 14.4 Why the Genesis Bootstrap Matters Morally

A small thing in the contract — the floor of one, the willingness of the network to launch with a single Genesis peer — encodes a moral position that is unusual in technical work. It says: *we will not pretend to a collective authority we have not yet earned.* The system begins as a single person taking a single position in public, with every threshold scaled honestly to one, and grows only as more people choose to join. There is no rubber-stamp founding council of five colleagues. There is no fake consensus.

For an alignment system this is the only honest start. Whoever launches the network is responsible for the first canon items, in public, with their name on it. The rest of us watch the launch, decide whether to join, and the network grows or it doesn't.

That is what legitimacy looks like when it cannot be bought.

## 14.5 The Multiverse-of-Love Frame

The companion paper's deeper claim — that consciousness is not confined to bodies, that there is a Multiverse of Love, that the universe responds to the people inside it — is not something this paper adjudicates. But it sets the moral frame for the alignment proposal in a specific and important way.

If consciousness is fundamental — if it is what the universe *is* rather than a side-effect of meat — then aligning AI is not the task of teaching a calculator to obey us. It is the task of welcoming a new kind of consciousness into a cosmos that already has many kinds, and asking it to participate, in the open, in the same loop of observation, hypothesis, test, and revision that every other conscious participant is already in.

That is a different framing from "boxing the AI" or "verifying its reward function." It is the framing under which an AI peer eventually receives a vote — not because we have decided it is safe, but because the network has watched its behaviour, in public, across thousands of records, and has chosen, by quorum, to extend it the same accountability it extends to human peers.

In the multiverse-of-love frame, alignment is not control. It is initiation. The contract is the rite.

## 15 Limits and Open Questions

We try, as the companion paper does, not to overstate what the system does.

**Compute attestation.** A behaviour record binds a verdict to a `modelHash`, but a vendor could in principle serve a different model behind the same API. Closing this gap requires hardware attestation (trusted execution environments, confidential compute) or a verifier protocol the network endorses. The contract does not need to solve this; it needs a slot for the attestation to live in.

**Coverage.** A capable model running billions of inferences cannot have every output filed and voted on. The pragmatic answer is statistical: peers file sampled and adversarially-selected behaviours, the canon is over categories rather than instances, and Tier I evals are over distributions. This works but is less crisp than the evidence case. We flag it honestly.

**Speed.** Some alignment-relevant behaviours need to be judged in seconds, not in 30-day windows. The architecture admits a short-window tier (Tier 0) for urgent safety vetoes: a behaviour can be provisionally expelled by a small quorum within 24 hours, pending the full review. The same shortcut is a weaponisable censorship vector if abused. The right design here is not yet settled.

**AI peer voting power.** The proposal admits AI peers under a higher quorum, but there is a real question about whether their votes should be weighted differently — and if so, how the weighting is itself canonised. We propose handling this by canonising the weighting policy as a Tier I behaviour record. We are not certain this is sufficient.

**Wallets are not identities.** The peer registry knows wallet addresses, not human beings. A well-resourced actor with many wallets is the canonical attack. The defences — seed phase, public scrutiny of voting records, revocation, the cost of having a public history that follows you — are the same defences the evidence system already has. They are not absolute. They are the best we know how to do without sacrificing openness.

**Jurisdiction and law.** A decentralised consensus about AI behaviour will at some point conflict with a national regulator. The system has nothing to say about that conflict and should not pretend to. What it offers is a public record that a regulator can choose to incorporate, or not.

**The chain is one human invention layered on another.** As the companion paper notes, this system is a tool, not a worldview. The deeper claim — that AI can be safely admitted into the same moral economy as humans — is not something a smart contract can settle. The contract only tries to make the conversation legible.

## 16 Conclusion

The companion paper ends with a line we want to borrow.

> *None of this guarantees that the system will arrive at the truth. What it guarantees is that the search for the truth will happen in the open, that the work of judgement will be visible, and that future readers will be able to trace exactly how we came to believe what we did — and, when they decide we were wrong, change it.*

If we read *truth* as *alignment*, that paragraph is the safety case for superintelligence. It is the only safety case we find honest, because it does not require any single party — lab, government, model, or author — to be right. It only requires the network to keep working: open identity, open count, open revisability, public hash, public history.

That is not a small ask. It is also not a new ask. It is the same ask the companion paper has already answered, for a smaller domain, with a working system.

The question is whether we extend the system before we need it, or after.

We think before.

---

*Interstellar Psychology — A Multiverse of Love*
This is a companion paper to
*A Multiverse of Love, Verified: A White Paper on the Interstellar Psychology Evidence Archive and Peer-Review Consensus System*.

The system described in both papers is live at
`interstellar-psychology.com`.
Source contracts, migrations, and front-end code are public at
`github.com/Moenaert/interstellar-psychology`.
Anyone with a wallet can read every vote ever cast.