

A Multiverse of Love, Verified

A White Paper on the Interstellar Psychology
Evidence Archive and Peer-Review Consensus System

Neo

Interstellar Psychology

May 16, 2026

Abstract. Interstellar Psychology is a meta-scientific field that bridges rigorous evidence with phenomena traditionally dismissed by mainstream science: psychedelics, telepathy, mindsight, remote viewing, out-of-body experience, non-human intelligence, multiverse signal, and the fractal nature of reality. This paper describes the system built to keep that field *honest*: a public evidence archive backed by an on-chain peer-review consensus mechanism. Every claim is filed, weighted by source quality, voted on by verified peers, and recorded in an immutable log that anyone can audit. Canon is never final — any peer can challenge a previously accepted claim, and the network can deprecate or reaffirm it. The result is a self-correcting record that takes the spirit of the scientific method — replication, scrutiny, openness, revisability — and gives it teeth that no individual, journal, or institution can quietly override.

Contents

1	Introduction: Why a New System Is Needed	2
2	What Counts as Peer Review	3
2.1	The Scientific Method, Briefly	3
2.2	What Peer Review Adds	3
2.3	What Traditional Peer Review Lacks	3
3	Goals of the System	4
4	System Architecture	4
4.1	Layer 1: The Chain (Truth)	4
4.2	Layer 2: The Cache (Speed)	4
4.3	Putting Them Together	5
5	The Subject Matter: Nine Pillars	5
6	Evidence: Tiers	5
7	The Seven-State Lifecycle	6
8	Who Can Vote: The Peer Registry	7
8.1	The Genesis Peer	7
8.2	The Seed Phase	7
8.3	Nomination and Endorsement	7

8.4	Revocation	7
8.5	Quorum Scaling, in One Table	8
9	Submitting Evidence	8
9.1	Stage One: Anyone Can File	8
9.2	Stage Two: A Peer Registers It On-Chain	8
10	Reviewing Evidence	8
11	Challenging Canon	9
12	Integrity: How We Know Nothing Has Been Quietly Edited	9
13	The Larger Goal	10
14	Comparison to Traditional Peer Review	11
15	Limits and Open Questions	11
16	Conclusion	11

1 Introduction: Why a New System Is Needed

Science advances by a simple loop: someone proposes a claim, others examine it, reproduce it, and either accept it into the shared record or push back against it. The institution we call *peer review* is the gate that decides what enters that record. It is a remarkable invention, and over the past three centuries it has produced most of what we know about the physical world.

It is also a system with two well-known weaknesses.

1. **It is private.** A handful of anonymous reviewers, chosen by an editor, decide what is publishable. The deliberation is invisible; the criteria are unwritten; the rejections do not travel with the paper.
2. **It is conservative.** A reviewer who has staked a career on one paradigm is the worst possible judge of evidence that contradicts it. Topics that fall outside the current consensus — regardless of the quality of the data — are routinely filtered out before publication.

This second weakness is the founding problem of Interstellar Psychology. A great deal of replicable, recorded, sometimes *government-declassified* evidence exists for phenomena that the mainstream is unwilling to host: twenty-three years of CIA remote-viewing research, congressional testimony on non-human intelligence, controlled studies of psychedelic healing, on-camera demonstrations of non-speaking autistic telepathy, cardiac-arrest patients describing operating rooms from the ceiling. The data is not in dispute. The *venue* is.

Interstellar Psychology builds that venue. It is a public archive of evidence, organised by topic, weighted by source quality, and curated by a network of verified peers who vote in the open and whose votes are written to a public ledger that no one can erase.

This paper describes how that system works, in plain language, and why we believe it is a fair upgrade to the traditional peer-review loop.

2 What Counts as Peer Review

Before describing what we built, we should be clear about what we are keeping and what we are changing.

2.1 The Scientific Method, Briefly

The scientific method is not a single procedure. It is a posture toward truth that has roughly four moves:

1. **Observation.** Notice something that asks for an explanation.
2. **Hypothesis.** Propose an explanation that could be wrong.
3. **Test.** Design an experiment or gather evidence that could contradict the hypothesis.
4. **Revision.** If the evidence contradicts the hypothesis, change the hypothesis. If it survives, accept it provisionally and keep testing.

The crucial property is the last one: every accepted claim remains *provisional*. Nothing in science is final; everything is one counter-example away from being revised. This is the property our system most wants to preserve.

2.2 What Peer Review Adds

Peer review is a social layer on top of the method. Before a claim enters the public record, other practitioners check the work:

- Are the methods sound?
- Is the data credible?
- Do the conclusions follow from the data?
- Is the work honest about its limitations?

If enough qualified reviewers say yes, the claim is published. If not, it is sent back or rejected. Peer review is the community’s way of saying “we have looked, and we collectively stand behind this.”

2.3 What Traditional Peer Review Lacks

Reviewers are anonymous. Editors have final say. The losing arguments are not preserved. There is no public count of *how many people* endorsed the claim or contested it. There is no way for a later generation to revisit an old paper and say “the network now thinks this is wrong” without writing a whole new paper.

Our system keeps the four moves of the method, keeps the social check, and adds three things traditional peer review does not have:

1. **Open identities.** Every reviewer is publicly named. Their endorsements and challenges follow them.
2. **Open tallies.** Every vote is visible, counted in real time, and stored permanently.
3. **Open revisability.** A canonized claim is not closed. Any verified peer can re-open it for the whole network to re-judge.

3 Goals of the System

We designed the Interstellar Psychology evidence system to satisfy five goals at once. These goals constrain everything that follows.

- 1. Truth is provisional.** No claim can be locked in forever. The network must always be able to revisit and revise.
- 2. Curators are accountable.** Whoever endorses an item must do so publicly, with a wallet that follows them across decisions.
- 3. The record is tamper-evident.** A claim cannot be silently edited after it has been accepted. Any later mutation must be detectable from outside the system.
- 4. The bar is calibrated to the evidence.** A peer-reviewed paper does not have to clear the same hurdle as a first-person account. The system must encode this distinction explicitly.
- 5. It must work from day one.** The network must be able to launch with a single founder and grow safely from there — without leaving a window where one person can rubber-stamp themselves into a fake quorum.

4 System Architecture

The system has two layers that work together. We separate them deliberately so that one of them can fail without corrupting the other.

4.1 Layer 1: The Chain (Truth)

A smart contract called `EvidenceConsensus`, deployed on the BNB Smart Chain, holds the canonical state of the network. It is the authoritative record. It tracks:

- Who is a verified peer.
- Who has been nominated to become a peer.
- Every piece of evidence submitted, along with its tier and a cryptographic fingerprint of its content.
- Every vote: approvals, rejections, challenges, defences.
- The current state of every claim (Submitted, Canon, Expelled, Lapsed, Contested, Deprecated, Reaffirmed).

Because it lives on a public blockchain, nothing on this layer can be retroactively rewritten. Every action is recorded, signed by a peer’s wallet, and visible to anyone with an internet connection.


4.2 Layer 2: The Cache (Speed)

A traditional database (Postgres, hosted on Supabase) holds the *readable* version of everything: full evidence text, summaries, search indexes, peer handles. This layer exists for one reason — speed. Reading from a blockchain is slow and expensive; reading from a database is instant.

The cache is a *follower*, not a source of truth. A background process (the “indexer”) continuously reads new events from the chain and updates the database to match. If the cache ever disagrees

with the chain, the chain wins. A daily integrity audit re-hashes every stored claim and compares it to the fingerprint recorded on-chain at submission time. Any mismatch raises a public alert.

4.3 Putting Them Together



A peer's vote is signed in their wallet, sent to the contract, and written to the chain. The indexer notices the new event and updates the cache. The website re-renders. The whole loop is typically visible within a minute.

5 The Subject Matter: Nine Pillars

Evidence is organised under nine topical pillars. They are the research domains that Interstellar Psychology investigates:

1	Music	Sound as the harmonic substrate that lets souls recognise each other.
2	Psychedelics	Compounds that can reproduce mystical experiences in a safe setting.
3	Telepathy	Non-speaking autistic individuals demonstrating mind-to-mind communication on camera, repeatedly.
4	Mindsight	Children trained to read text and identify colours while fully blindfolded.
5	Remote Viewing	Twenty-three years of declassified CIA research into non-local sight.
6	Out of Body	Cardiac arrest survivors describing operating rooms from the ceiling.
7	Non-Human Intelligence	From AAWSAP through congressional testimony — the question is no longer whether.
8	Multiverse	Synchronicity as evidence. The universe signalling that it sees you.
9	Infinity	Self-similar, scale-invariant, endless. The fractal thumbprint of God.

These nine domains are not the claim of the system. They are the *question*. The claims live inside them.

6 Evidence: Tiers


Not all evidence is equal. A declassified intelligence document and a single first-person testimony both belong in the record — but they cannot carry the same weight. The system encodes this with three tiers:

Tier	Type	Example
I	Peer-reviewed or declassified	Published scientific paper; CIA STARGATE files.
II	Documented or institutional	Books, podcasts, documentaries, official testimony.
III	Testimony or first-person	Personal experience reports; eyewitness accounts.

Tier I is the most credentialled and therefore requires the *highest* percentage of peers to canonize *and* the *highest* percentage to remove. It is harder to add and harder to lose. Tier III is the most permeable in both directions.

7 The Seven-State Lifecycle

Every piece of evidence travels through a finite sequence of states. The contract enforces the transitions; nothing else can change them.



Submitted. A claim has been filed and is awaiting peer review. There is a 30-day window during which peers may vote approve or reject. If neither side reaches a quorum, the claim *lapses*.

Canon. Enough peers approved the claim. It enters the public archive and is treated as accepted by the network.

Expelled. Enough peers rejected the claim. It is recorded as not accepted and is not displayed in the main archive.

Lapsed. The 30-day window passed without enough votes either way. The claim is set aside for lack of consensus rather than for being wrong.

Contested. A verified peer has formally challenged a canon claim. The piece is flagged in the archive, and a 21-day window opens for the network to vote on whether to deprecate it or defend it.

Deprecated. Enough peers supported the challenge. The claim is marked as no longer canon. The original record is preserved — it is not deleted — so that the history of the disagreement remains visible.

Reaffirmed. The challenge did not reach the deprecation threshold within the window. The claim is restored to canon, with a visible note that it survived a challenge. A reaffirmed claim can itself be challenged again later.

8 Who Can Vote: The Peer Registry

A claim's fate depends entirely on who counts as a peer at the moment of the vote. This is the most carefully designed part of the system.

8.1 The Genesis Peer

The network launches with a single founding peer, called the *Genesis* peer. At launch, the Genesis peer is the entire quorum. This is not a bug — it is the only honest starting condition. A new network has not yet earned the right to claim collective authority, and pretending otherwise (by, say, requiring five votes before anything can canon) creates a deadlock at day one.

To make Genesis-only operation safe, every threshold has a *floor of 1*. One peer is enough to approve, expel, or revoke, *but only while the network is still genuinely a network of one*. The numbers scale up automatically as new peers join.

8.2 The Seed Phase

There is one obvious abuse path: the Genesis peer could nominate a hundred fake wallets controlled by themselves and “verify” them all. To close this window, public nominations are locked until the active peer count crosses a configured threshold (called the *seed phase K*). During the seed phase, the contract's owner adds peers directly. Once enough independent peers exist that they can no longer be silently outvoted by sockpuppets, the seed phase ends and the public nominate–endorse flow opens.

8.3 Nomination and Endorsement

After the seed phase, any verified peer can nominate a new candidate by submitting the candidate's wallet address and a display name. Other verified peers then endorse the nominee. When endorsements reach the nomination threshold, the nominee is automatically promoted into the active peer list.

The nomination threshold scales with the network: roughly one third of the active peer count, rounded up, and capped at nine to keep onboarding fast in mature networks.

8.4 Revocation

A peer who acts in bad faith — voting maliciously, spamming the challenge system, repeatedly endorsing low-quality submissions — can be removed. Any peer can motion to revoke another peer, and others vote on the motion. When votes reach the revocation threshold (roughly half the network), the peer is automatically removed.

A peer cannot motion to revoke themselves, and the network can never vote itself down to zero — the contract enforces that at least one active peer must always remain.

8.5 Quorum Scaling, in One Table

Thresholds as a function of the active peer count n :	
Action	Threshold (peers required)
Canonize a Tier I submission	$\lceil 0.45 n \rceil$, minimum 1
Canonize a Tier II submission	$\lceil 0.35 n \rceil$, minimum 1
Canonize a Tier III submission	$\lceil 0.30 n \rceil$, minimum 1
Expel a pending submission	$\lceil 0.25 n \rceil$, minimum 1
Deprecate a Tier I canon item	$\lceil 0.65 n \rceil$, minimum 1
Deprecate a Tier II canon item	$\lceil 0.60 n \rceil$, minimum 1
Deprecate a Tier III canon item	$\lceil 0.55 n \rceil$, minimum 1
Verify a peer nominee	$\min(9, \lceil n/3 \rceil)$, minimum 1
Revoke a peer	$\lceil n/2 \rceil$, minimum 1

Notice the asymmetry. Tier I requires the most peers to enter (the bar is high) and the most peers to deprecate (it is also defended strongly once accepted). The system encourages care on the way in *and* on the way out.

9 Submitting Evidence

The submission flow exists in two stages so that ordinary readers can contribute without needing a wallet, while verified peers handle the on-chain step.

9.1 Stage One: Anyone Can File

A visitor opens the Evidence page, clicks *Submit evidence*, and fills in a form: pillar, type, tier, title, source, year, a short excerpt, a link, and tags. The submission is stored as a *pending* row in the cache. It does not yet appear in the public archive.

To prevent abuse, anonymous submissions are throttled at five per IP address per hour, enforced atomically inside the database.

9.2 Stage Two: A Peer Registers It On-Chain

Verified peers see all pending submissions in their queue. To bring one into the formal review process, a peer:

1. Recomputes a cryptographic fingerprint (`keccak256`) of the submission's canonical content.
2. Calls `submitEvidence` on the contract with the submission's identifier, its tier, and the fingerprint.
3. Signs the transaction with their wallet.

The fingerprint binds the on-chain record to a specific version of the text. From this point on, any silent edit to the cached content can be detected by re-hashing and comparing to the on-chain value. This is what gives the cache its tamper-evidence.

10 Reviewing Evidence

Once on-chain, the submission enters the public review queue. Verified peers see it in their Peer Review dashboard and decide whether to approve or reject. There are several properties worth

noting:

- **Doubly signed.** Each vote produces two artefacts in parallel. First, an EIP-712 *attestation* — a structured, human-readable message signed inside the peer’s wallet that names the evidence, the phase, and the verdict. This signature is stored off-chain in the public attestation log. Second, the same peer submits an on-chain `castReviewVote` transaction, which is what the contract actually counts. The attestation gives a readable record of *why* a peer voted; the on-chain transaction gives the unforgeable count.
- **Final.** A peer cannot change their vote once cast.
- **Counted live.** The cache updates within a minute of the on-chain event. Anyone viewing the page sees the same running tally.
- **Batched.** Peers can vote on many items in one transaction (`castReviewVoteBatch`), capped at 50 items per call, to keep gas costs low.
- **Time-limited.** Submissions have a 30-day review window. After that the item lapses if no threshold was reached. Any peer can call `markLapsed` once the window closes; the function reverts if the window is still open.

11 Challenging Canon

This is the part of the system that most distinguishes it from traditional peer review.

A canon item is not closed. Any verified peer can challenge it at any time, with two safeguards:

- A peer who has just opened a challenge must wait seven days before opening another. This stops a single peer from weaponising the challenge button to spam the network.
- The challenge must state grounds — a written explanation of what specifically is wrong, misleading, or no longer supported. The grounds are stored alongside the challenge so that voters know what they are voting on.

Opening a challenge moves the claim to *Contested* and starts a 21-day window. During the window, peers vote either to *challenge* (deprecate) or *defend* (reaffirm). Resolution is deterministic:

- If challenge votes reach the deprecation threshold before the window expires, the claim is immediately deprecated.
- If the window expires without reaching the threshold, the claim is automatically reaffirmed. Silence defends canon.

A reaffirmed claim carries a visible badge showing that it survived a challenge, along with the date. It can be challenged again later — the system never closes the loop on revisability. Each new challenge cycle starts from a clean slate of votes; previous cycles are preserved as part of the historical record.

12 Integrity: How We Know Nothing Has Been Quietly Edited

The most subtle failure mode of any archive is that someone changes the content of an old entry without anyone noticing. We defend against this with several independent mechanisms.

Content fingerprints. At submission time, the contract records a `keccak256` hash of the canonical content. This hash is part of the on-chain event log forever.

Cross-check at registration. When a peer submits an item on-chain, the edge function verifies that the recomputed hash matches the one in the transaction. A submission whose stored text does not match the hash on the transaction is rejected.

Daily audit. A scheduled job re-hashes every canon item in the cache and compares the result to the on-chain hash. Any mismatch raises a *tamper alert* that is visible on the public Peer Review dashboard. Alerts are not silent.

Heartbeats. Every scheduled job in the system writes a heartbeat after each run. A stale heartbeat surfaces on the dashboard so that operators *and the public* can see immediately if the indexer or audit job has stopped. A silent system cannot quietly fall behind.

Two public logs. The Peer Review dashboard exposes two read-only surfaces that anyone, peer or not, can browse: the *attestation log* (every EIP-712 signature ever produced, verifiable against the signer’s wallet) and the *chain log* (every on-chain event emitted by the contract, linked to the matching BscScan transaction). Together they make the deliberation behind every state change inspectable end-to-end.

Two-step ownership transfer. The contract owner cannot hand over administrative control in a single call. Transfer is split into `proposeOwner` and `acceptOwnership`: the proposed address must itself sign the acceptance, which rules out typos and contracts that cannot act. The transfer can also be aborted before acceptance with `cancelOwnershipTransfer`. This is a small detail; it removes a large failure mode.

13 The Larger Goal

A neutral observer might ask: why bother? Why build an entire on-chain voting system to curate evidence that mainstream journals would mostly decline to print anyway?

The answer is that the cost of *not* building it is high.

1. **Real phenomena go undocumented.** When credible researchers know their topic will not survive editorial review, they stop writing about it. The evidence does not disappear; it just stops being archived in a form that can be cited. A generation later, no one can find it.
2. **Bad evidence wins by default.** If the only people publishing on a topic are those who do not face institutional gatekeeping, the public record on that topic skews toward the least credentialed contributors. The good work and the crank work end up shelved next to each other.
3. **We forget that knowledge is revisable.** The longer an institution polices what counts as a fact, the easier it is to believe that the current consensus is the final answer. It never is.

The Interstellar Psychology archive is an attempt to give credible work on these topics a venue with a clear standard: open identity, open count, open revisability, public hash, public history. We are not asking anyone to take a single claim on faith. We are publishing *the work of judging the claim*, so that the judgement itself can be examined.

14 Comparison to Traditional Peer Review

	Traditional journal	Interstellar Consensus
Reviewer identity	Anonymous	Public wallet address
Number of reviewers	Typically 2–3	Scales with network ($\geq 30\%$)
Vote visibility	Hidden	Public, real-time
Acceptance criteria	Editor’s discretion	Algorithmic threshold
Record after publication	Static	Re-openable indefinitely
Retraction process	Editor-driven, opaque	Network vote, on-chain
Tamper-evidence	None	Content hash + daily audit
History of disagreement	Lost	Preserved
Cost to author	Often paid	Submission is free
Cost to reader	Often paywalled	Free to read forever

The traditional journal optimises for a clean public-facing record at the cost of erasing the deliberation behind it. We make the opposite trade: we keep the deliberation, every count, every challenge, every reaffirmation, on a public log. The record is messier, but the messiness is the point. It is what allows the record to correct itself.

15 Limits and Open Questions

We try not to overstate what the system does.

It does not adjudicate truth. The contract counts votes. It does not check whether the underlying paper is correct. It cannot. What it guarantees is that the count is honest and that no one has secretly edited the record afterwards.

It is only as good as its peers. A small or homogeneous peer set will produce a small or homogeneous canon. Growing the network — and growing it with people who disagree with each other in good faith — is the long project.

Wallets are not identities. The peer registry knows wallet addresses, not human beings. A determined attacker with enough resources can run multiple wallets. The seed phase, the revocation flow, and the public scrutiny of each peer’s voting record are the defences against this, but none of them is absolute.

Gas is not free. On-chain actions cost a small amount of BNB. We mitigate this with batched voting and by keeping the most expensive reads cached off-chain. We do not eliminate it.

The chain is one human invention layered on another. This system is a tool, not a worldview. The deeper claim — that there is a Multiverse of Love, that consciousness is not confined to bodies, that the universe responds to the people inside it — is not something a smart contract can settle. The system only tries to make the conversation legible.

16 Conclusion

Interstellar Psychology takes the four moves of the scientific method — observe, hypothesise, test, revise — and gives them an infrastructure that the original method never had:

- A public, append-only log of every claim.
- Public identities for everyone who endorses or challenges.
- Quorums that scale with the network and differ by tier so that the bar for entry matches the strength of the source.
- A revisability mechanism that means no canon item is ever truly closed.
- Tamper-evidence so that the record cannot be quietly rewritten.

None of this guarantees that the system will arrive at the truth. What it guarantees is that the search for the truth will happen in the open, that the work of judgement will be visible, and that future readers will be able to trace exactly how we came to believe what we did — and, when they decide we were wrong, change it.

That is enough. The rest is the work.

Interstellar Psychology — A Multiverse of Love

The system described in this paper is live at

interstellar-psychology.com.

Source contracts, migrations, and front-end code are public at

github.com/Moenaert/interstellar-psychology.

Anyone with a wallet can read every vote ever cast.